

# [320] Parallelism

Department of Computer Sciences  
University of Wisconsin-Madison

# Parallelism: doing multiple things at once

**Other Terms Today:** process, thread, instruction pointer,  
state (running, ready, blocked), CPU, GPU, core

## Outline:

- Mental Model
- Two problems
- Parallelism: Thread, Process, GPU

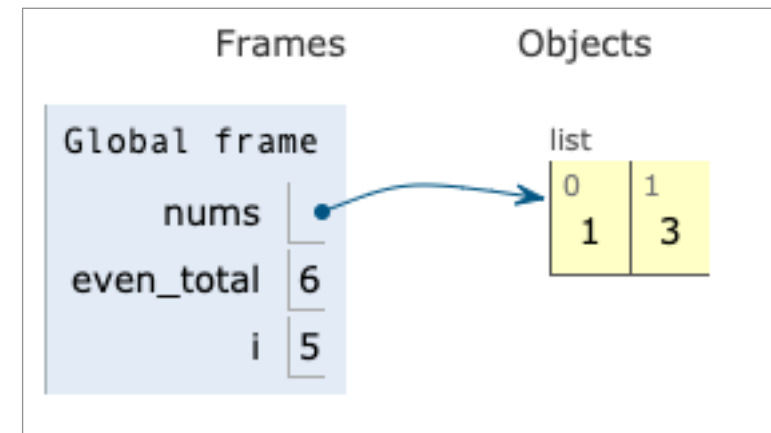
Mental Model: Tasks and Cores

# One Python Program Running


Code

```
1 nums = []
2 even_total = 0
3 for i in range(10):
4     if i % 2 == 0:
5         even_total += i
6     else:
7         nums.append(i)
8 print(i)
```


Data




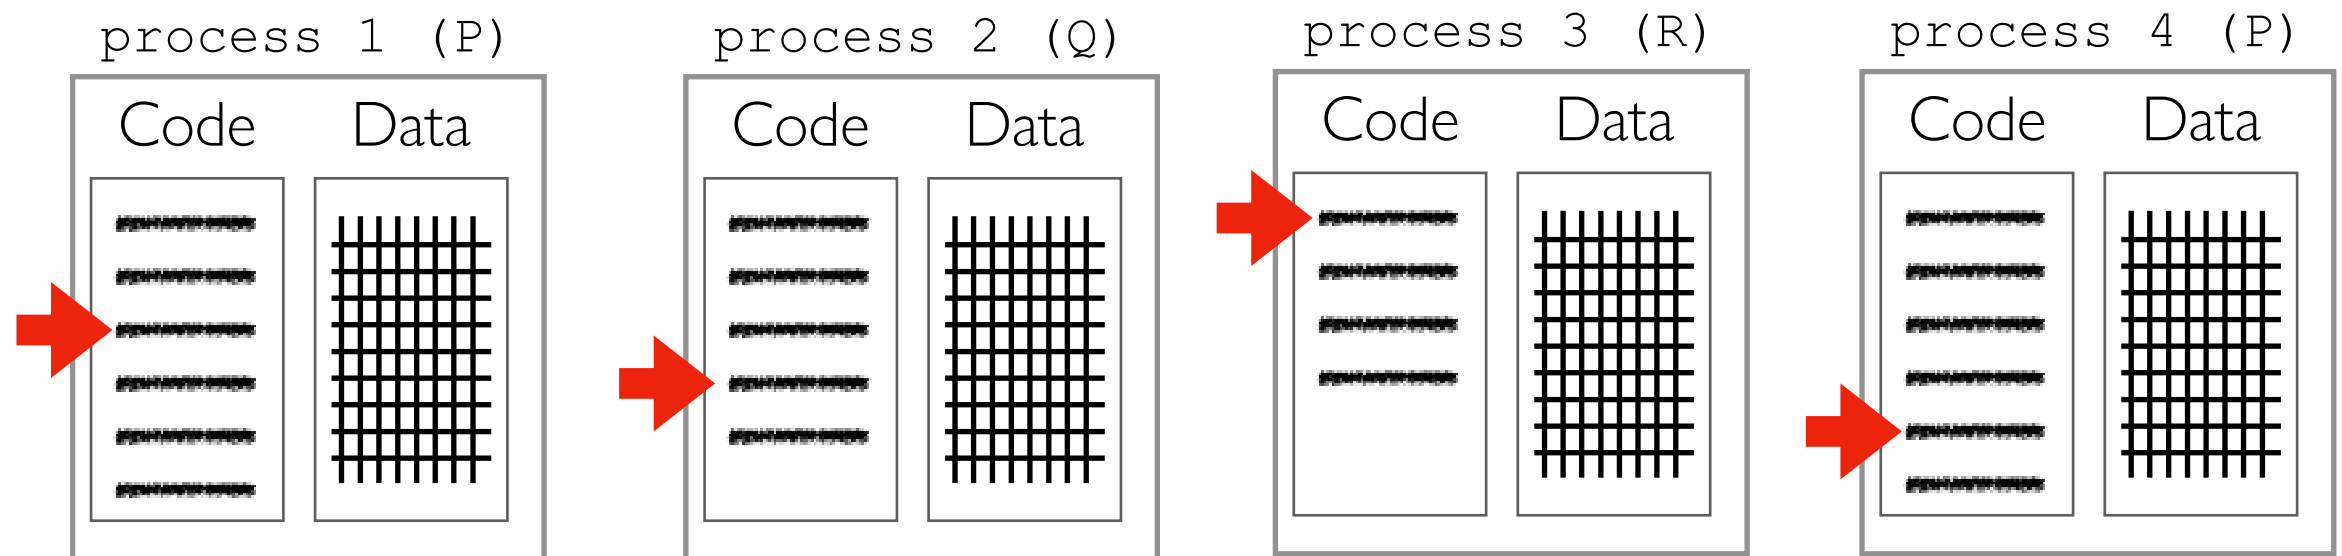
what is currently being done



instruction pointer  
(also called "program counter")

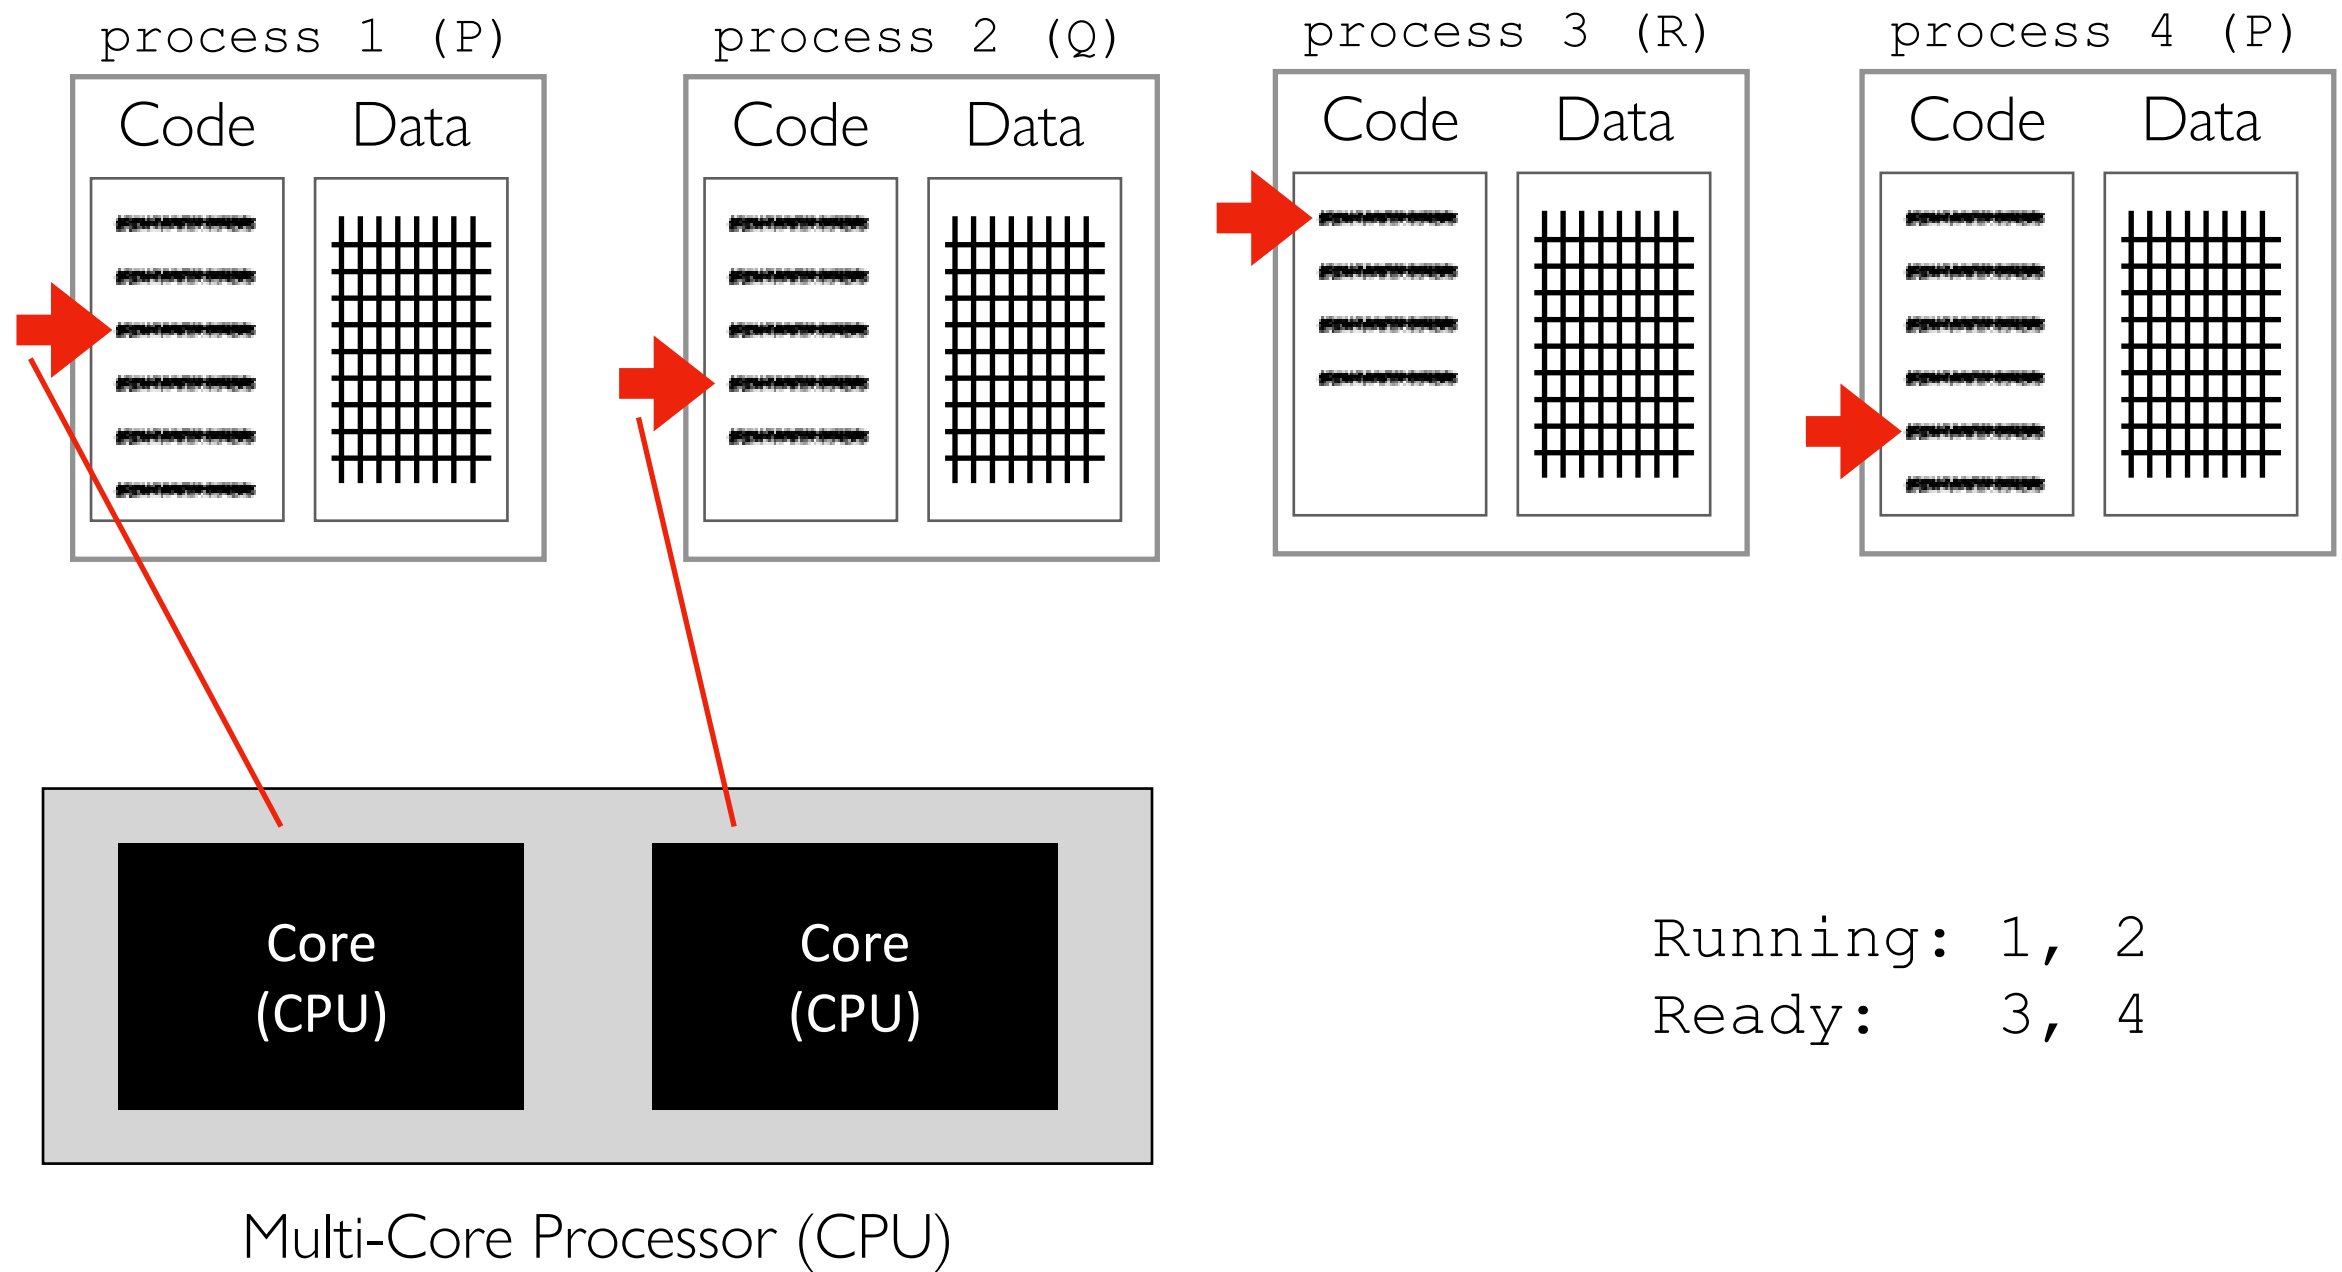


instruction pointer belongs to a *thread* within the process

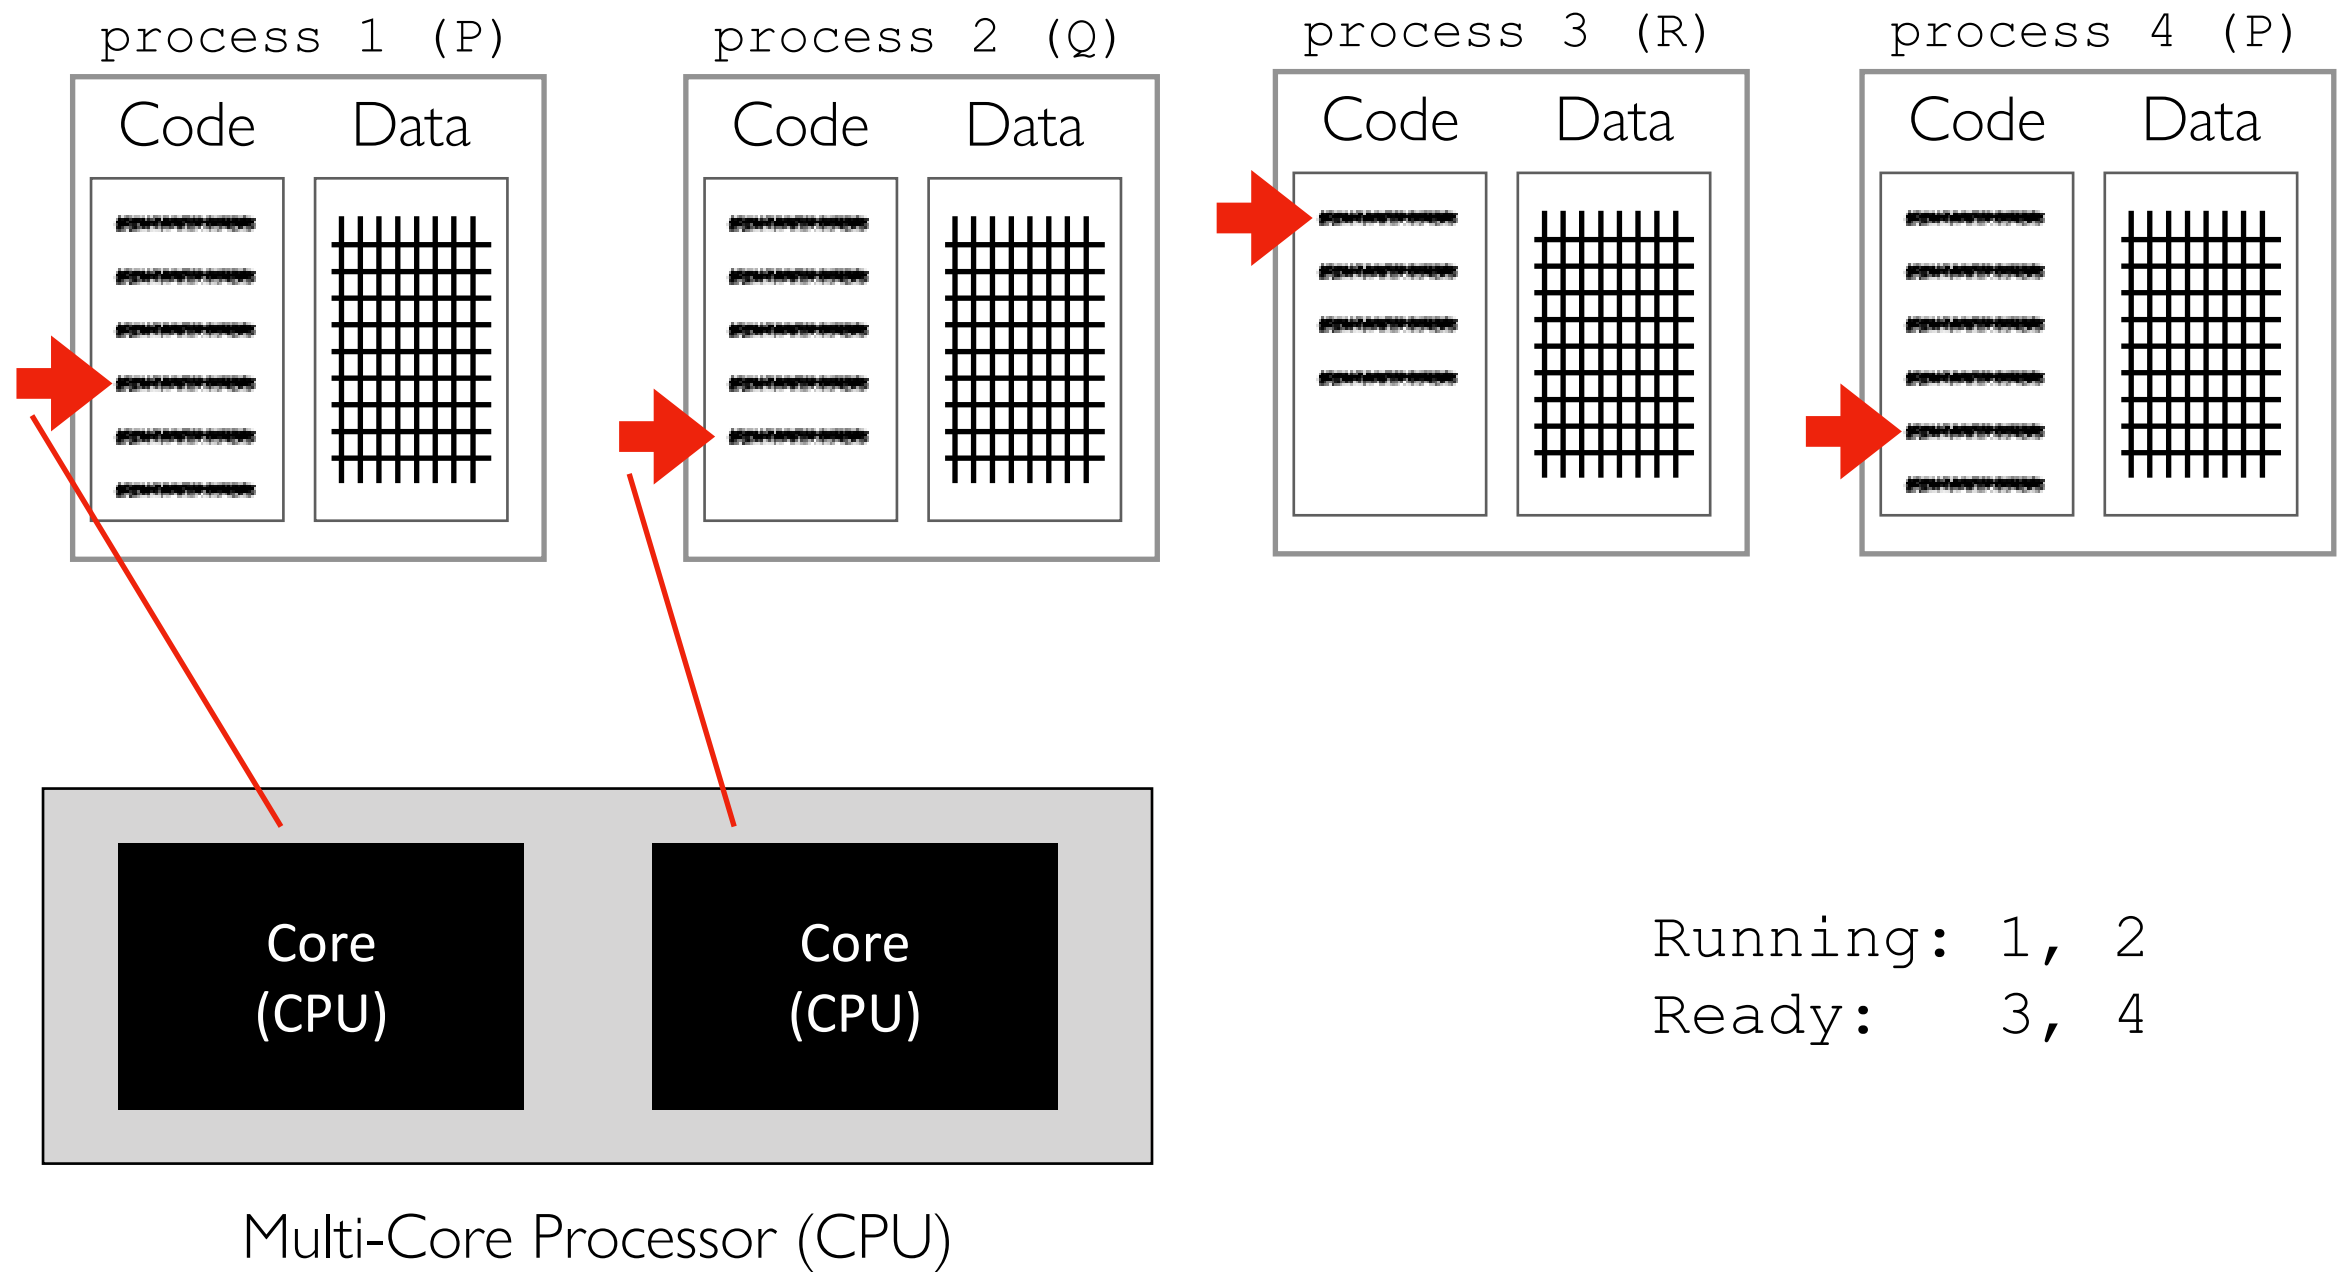


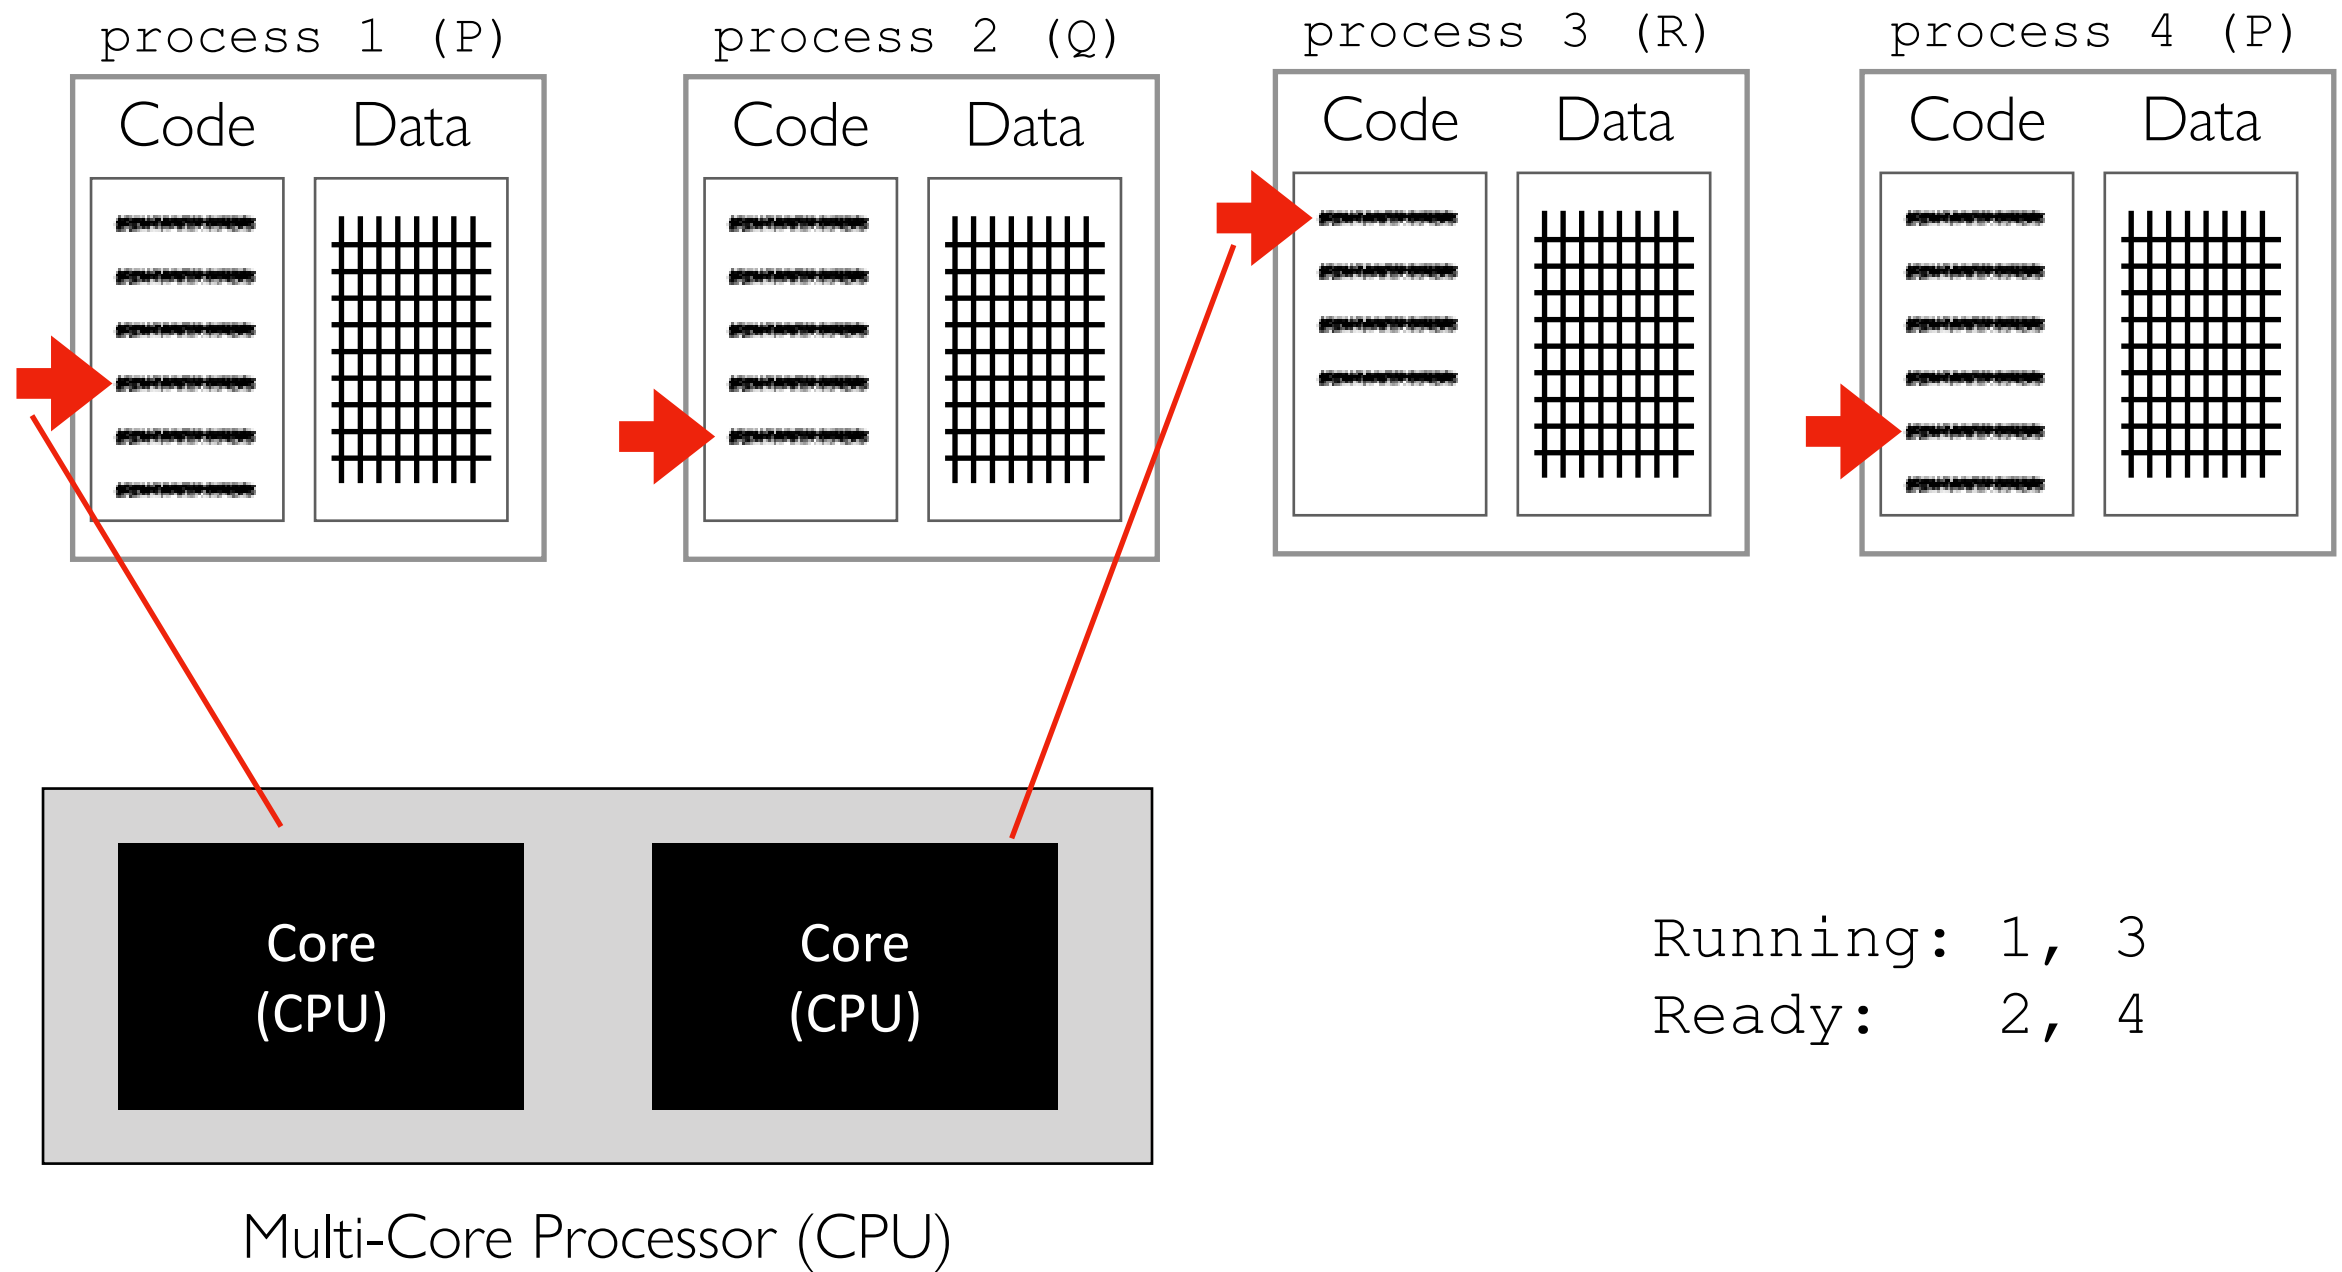
Multi-Core Processor (CPU)

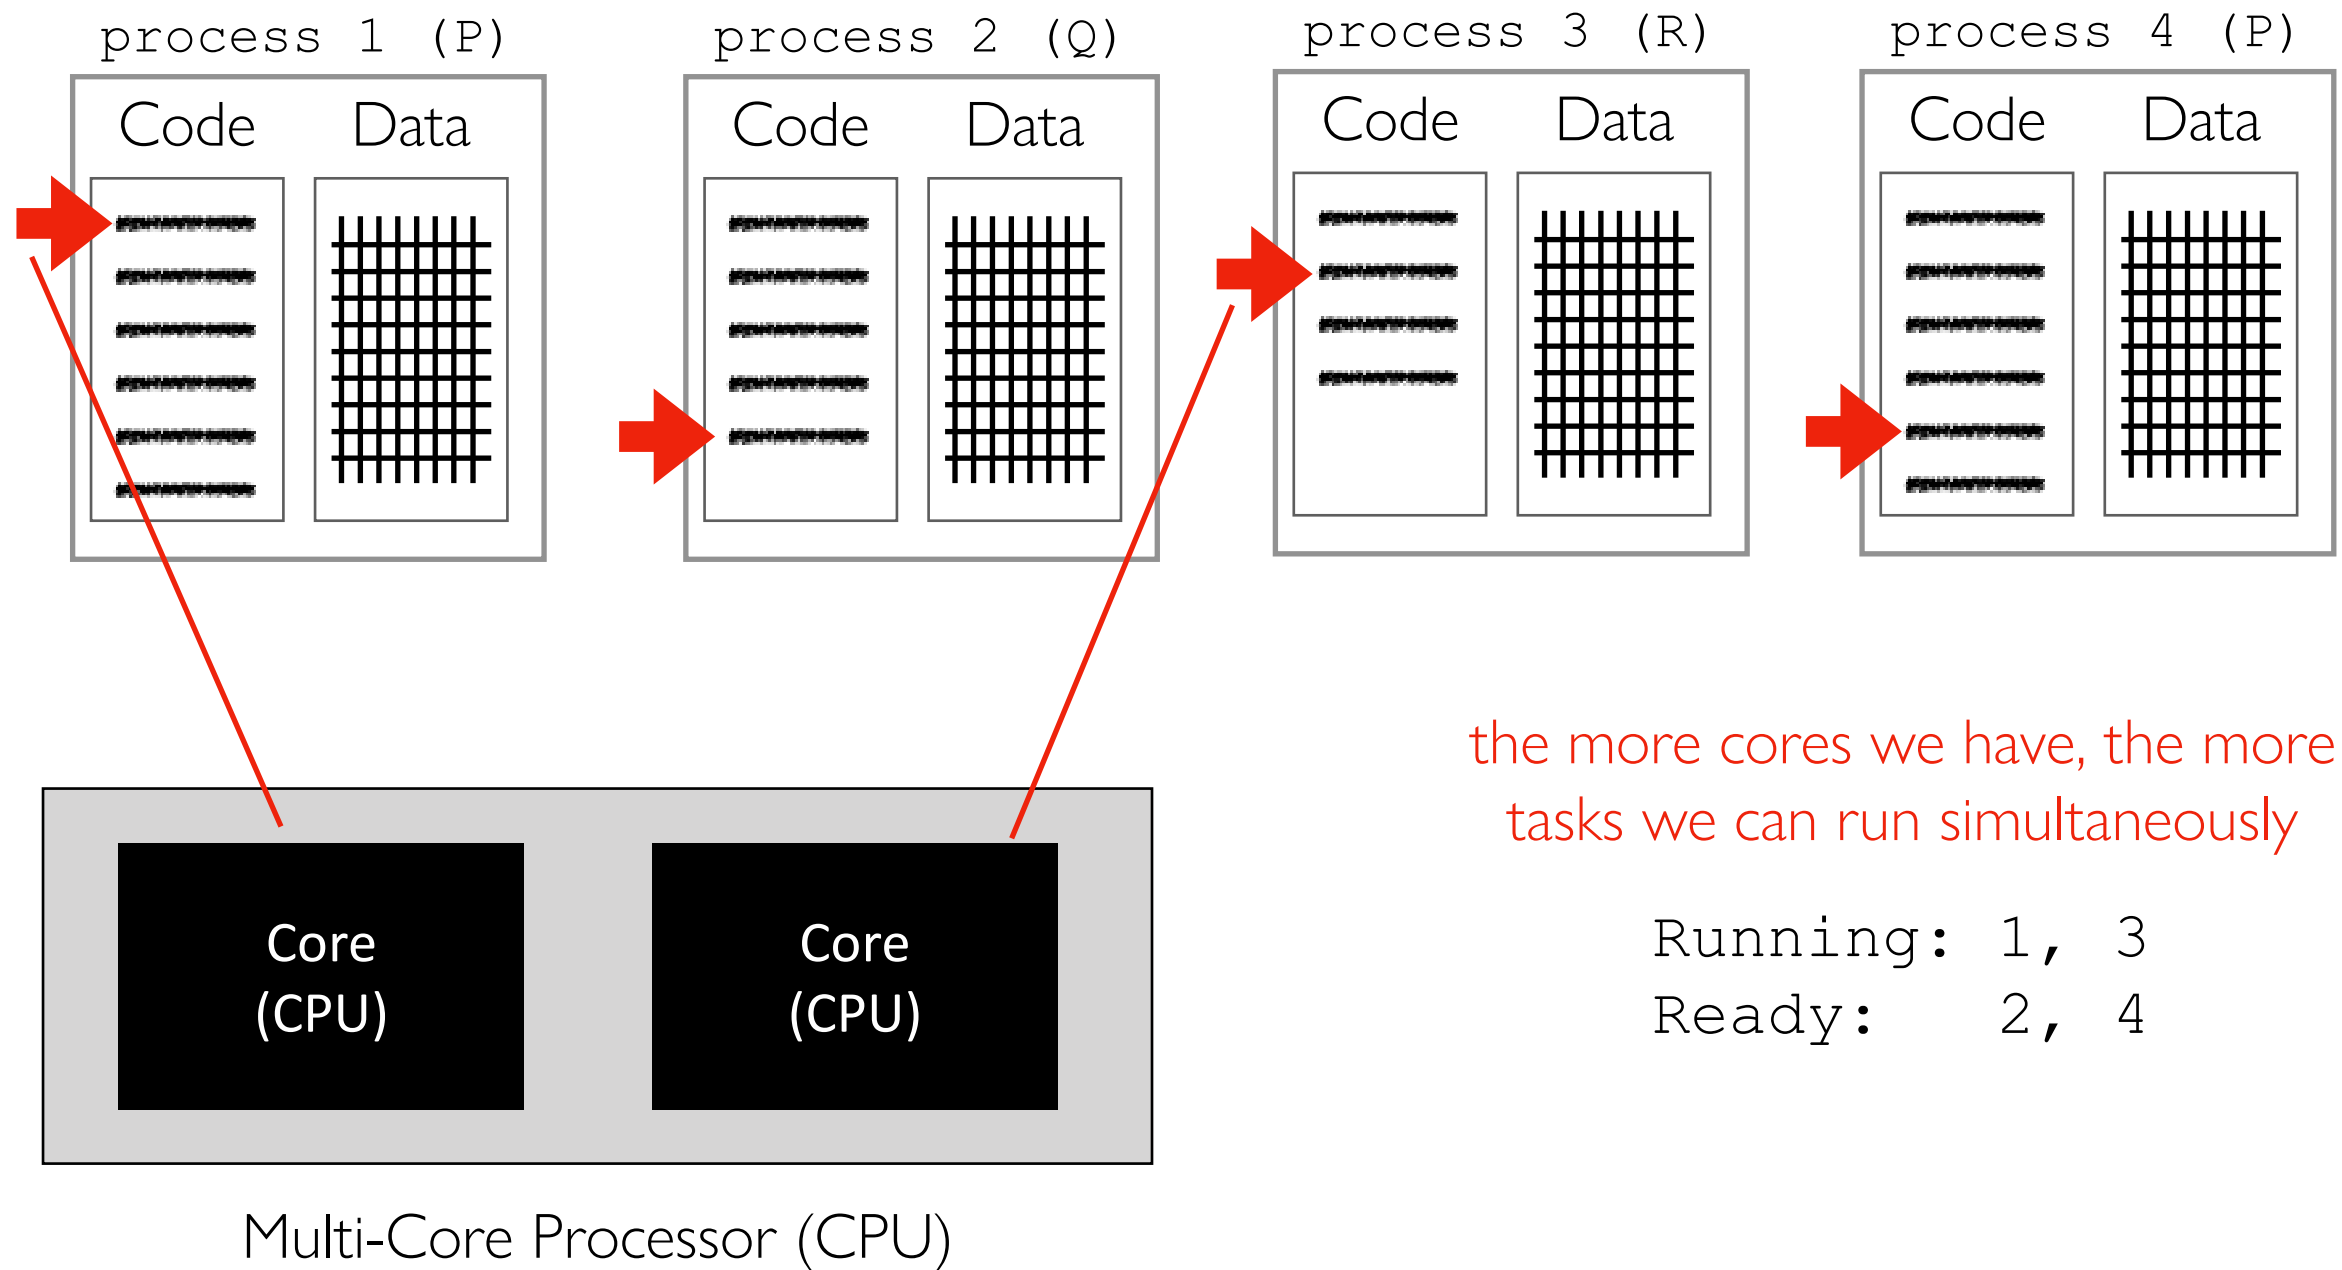






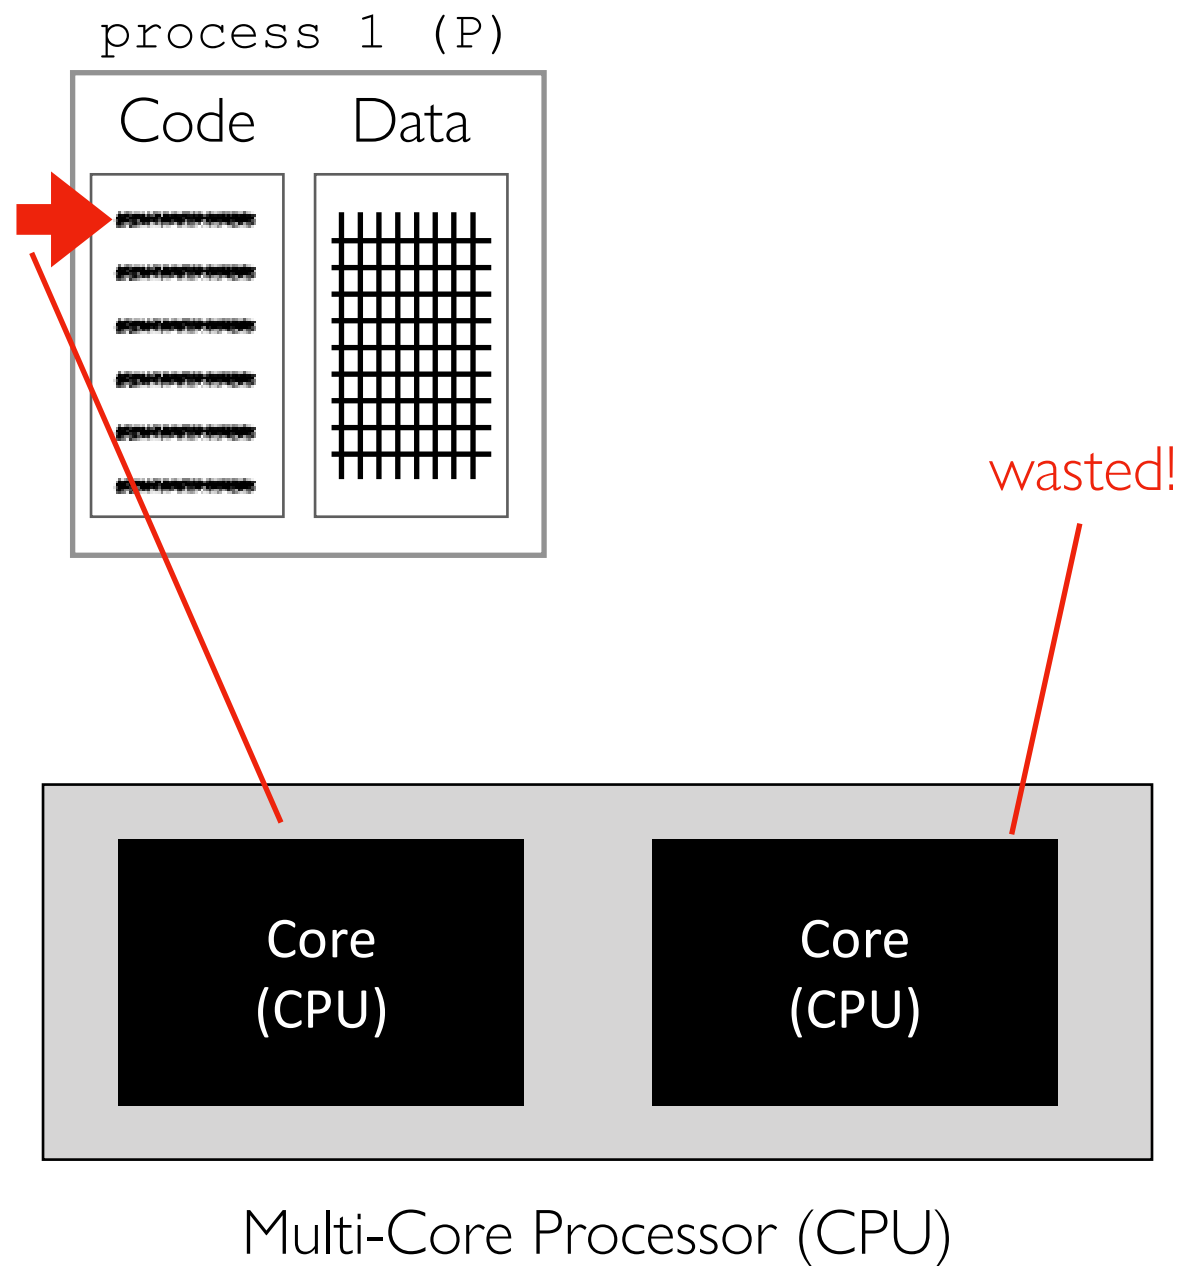






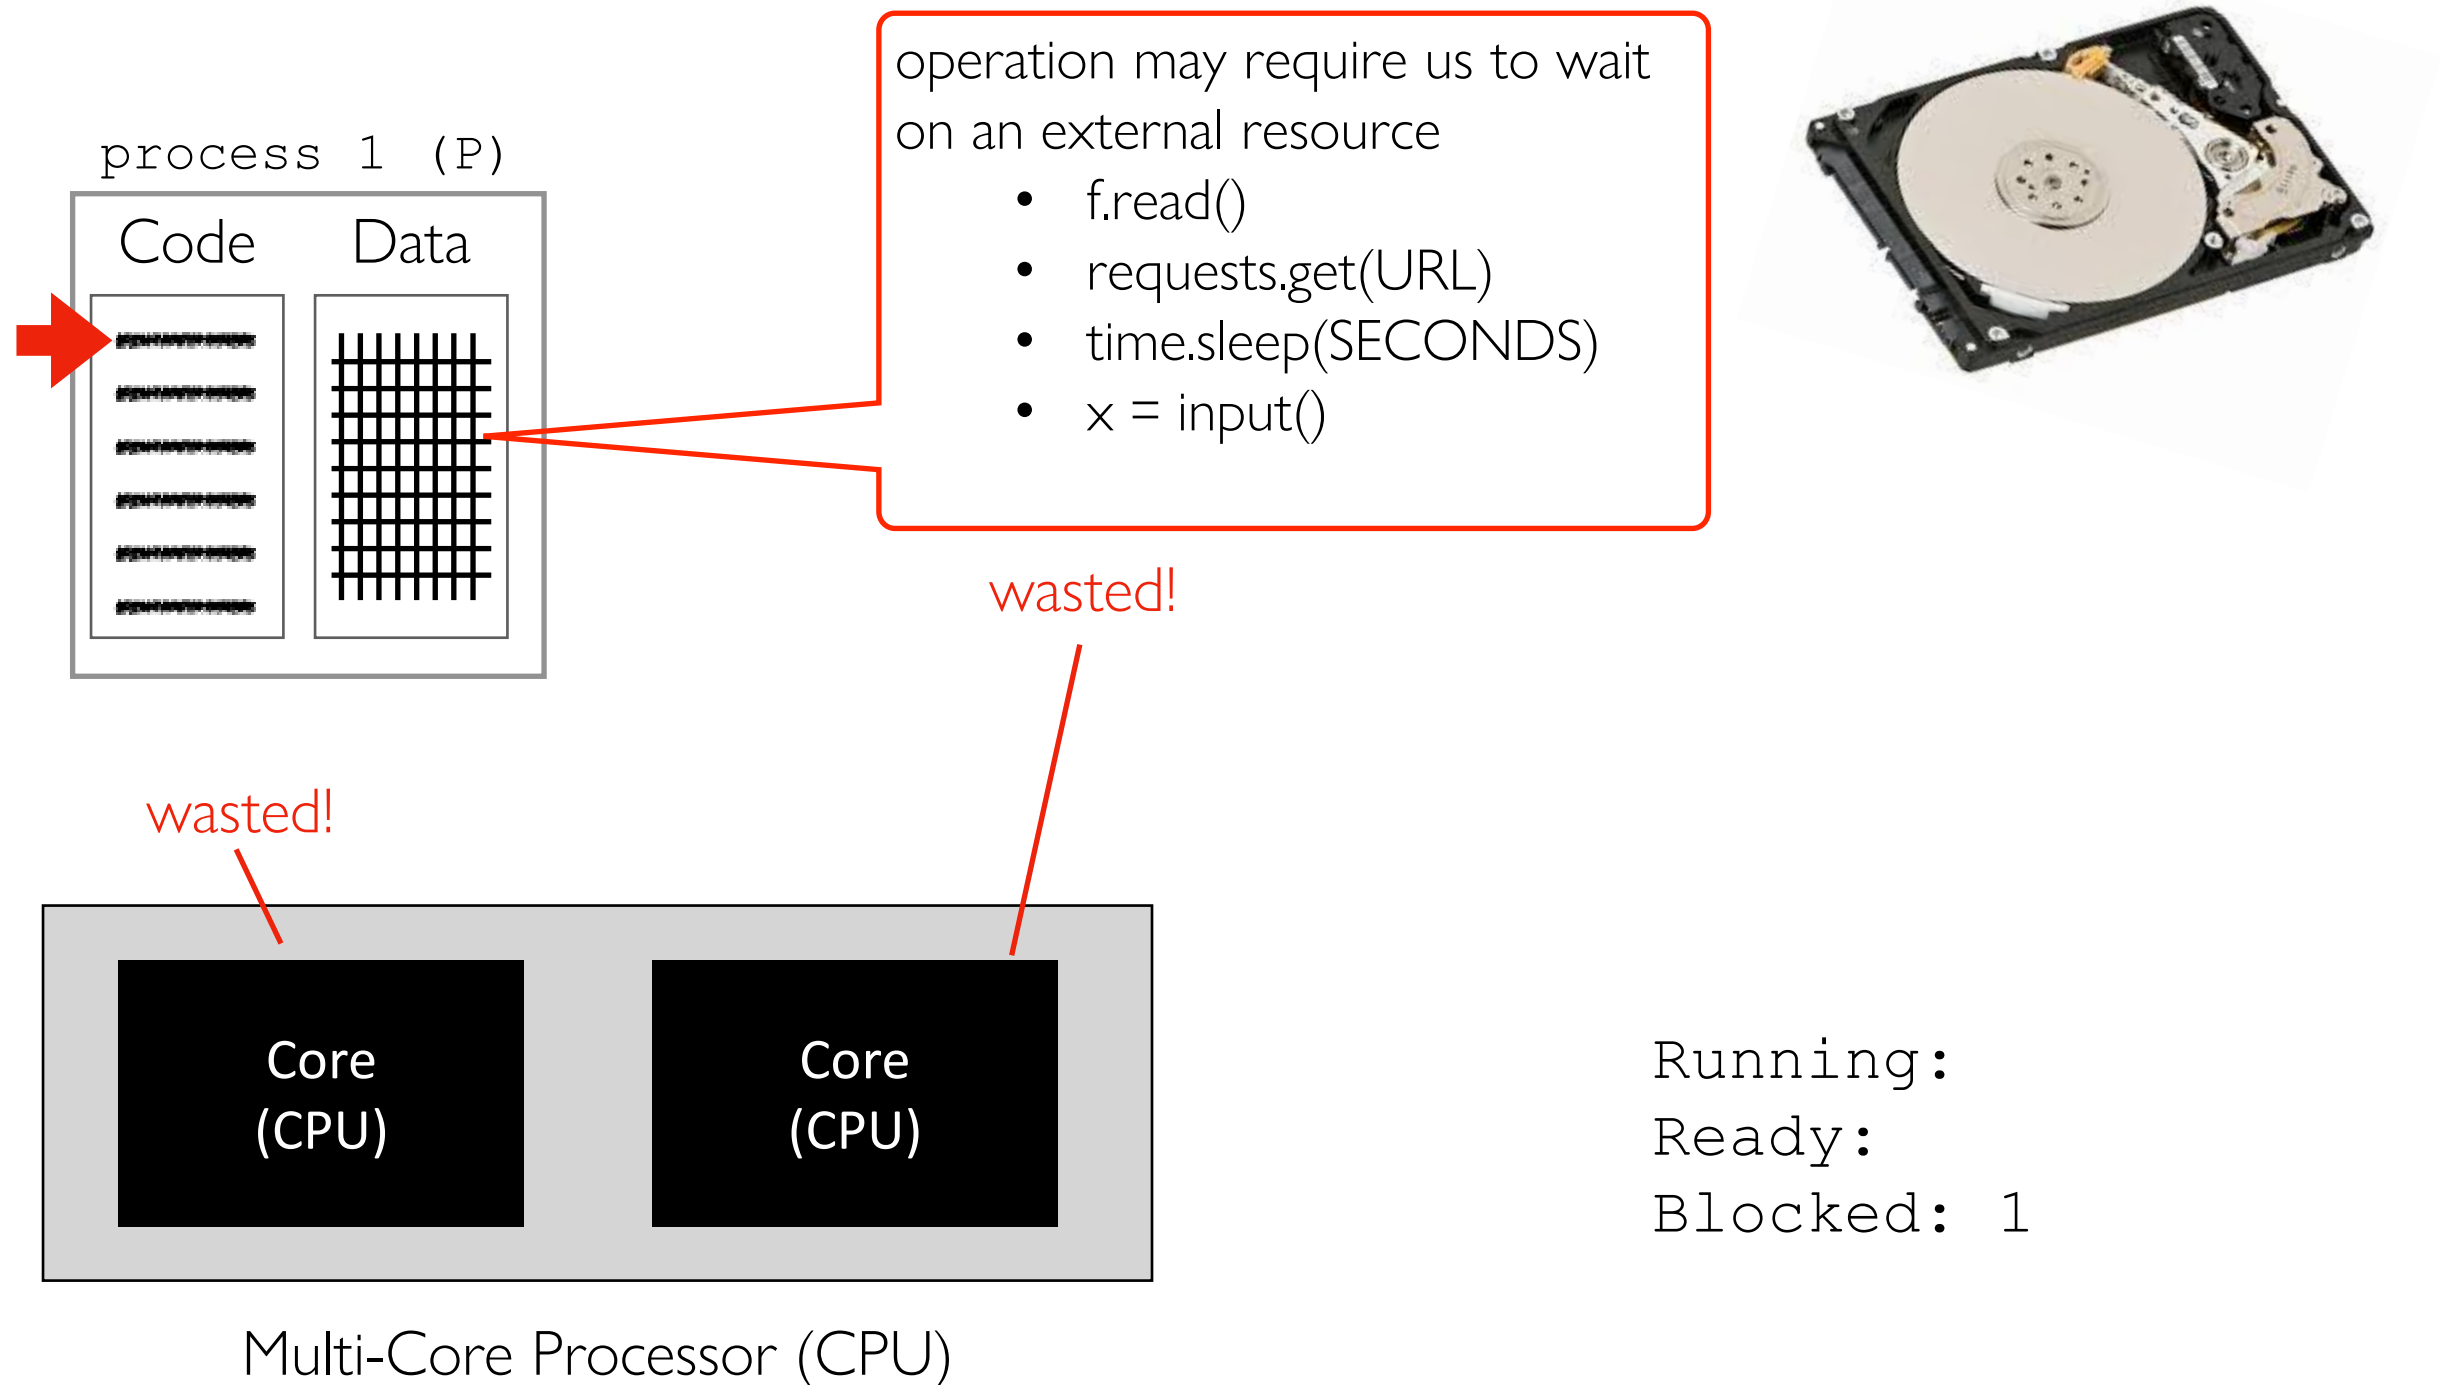
# Wasted Compute Resources: Two Problems

# Problem 1: not enough distinct tasks to utilize all cores



Running: 1  
Ready:


## Problem 2: some operations requires waiting (task is "blocked")




# Solution: Parallelism

- |   |                           |   |
|---|---------------------------|---|
| 1 | thread-level parallelism  | very complicated, not covered in detail |
| 2 | process-level parallelism | covered in CS 320                       |
| 3 | GPU parallelism           |   |

## (1) Thread-level Parallelism



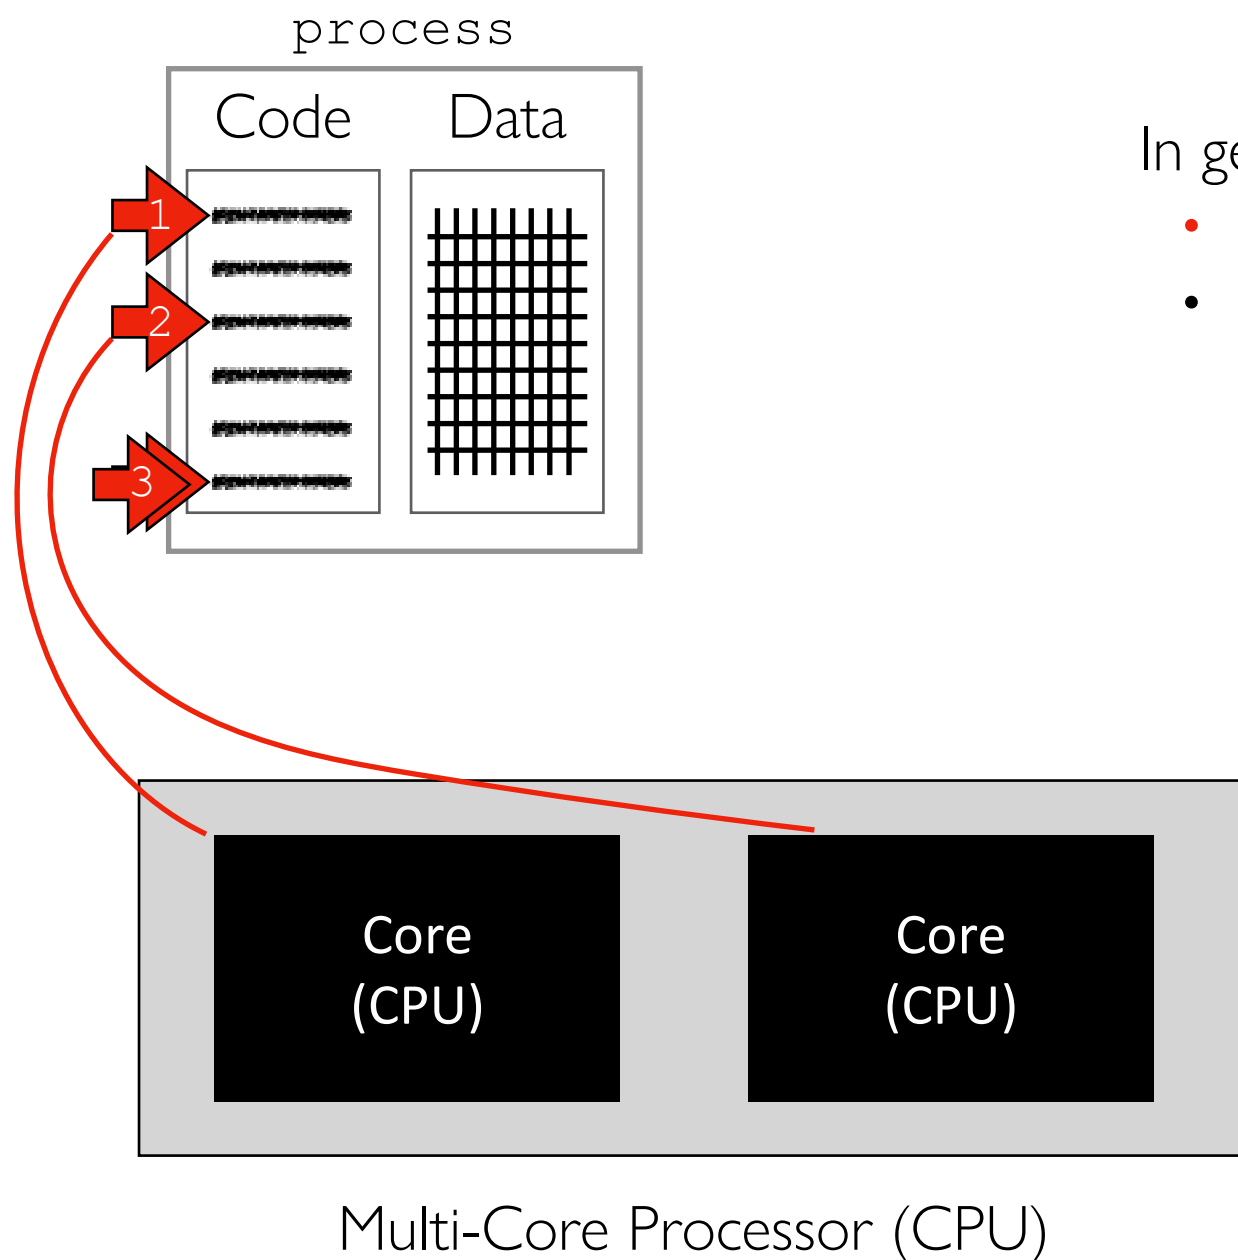
Threads give us multiple instruction pointers in a process, allowing us to execute multiple parts of the code, at the same time!



Multi-Core Processor (CPU)



# (1) Thread-level Parallelism

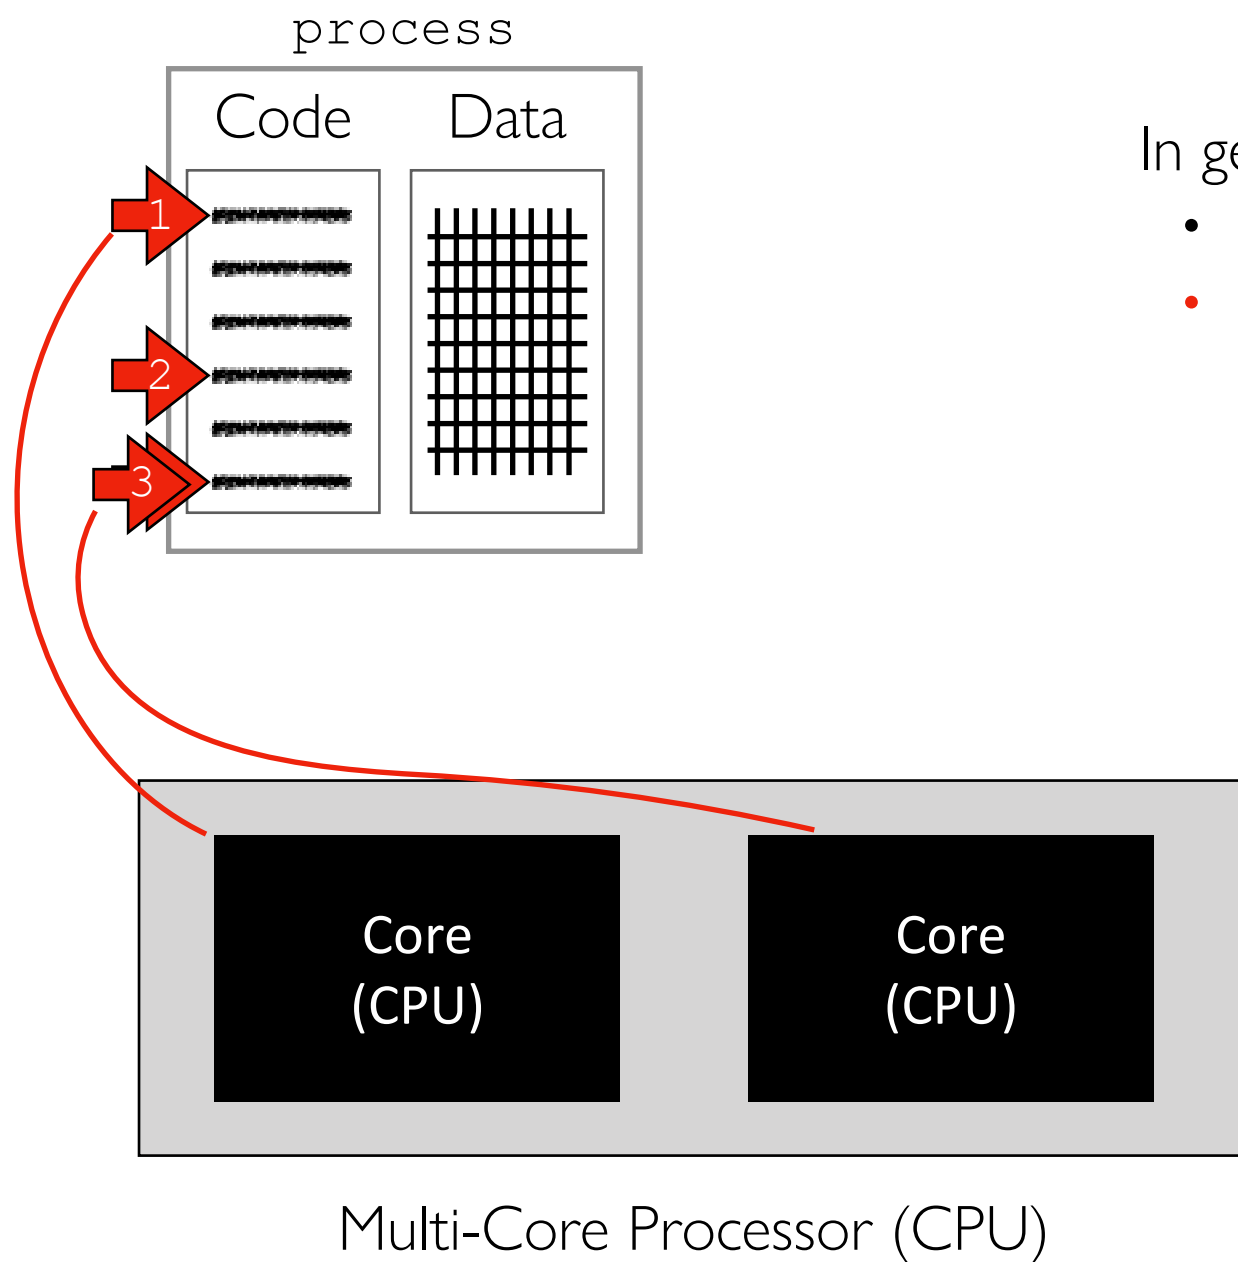


In general, threads help:

- use multiple cores
- do useful work when threads are blocking

Running: 1, 2  
Ready: 3, 4  
Blocked:

# (1) Thread-level Parallelism

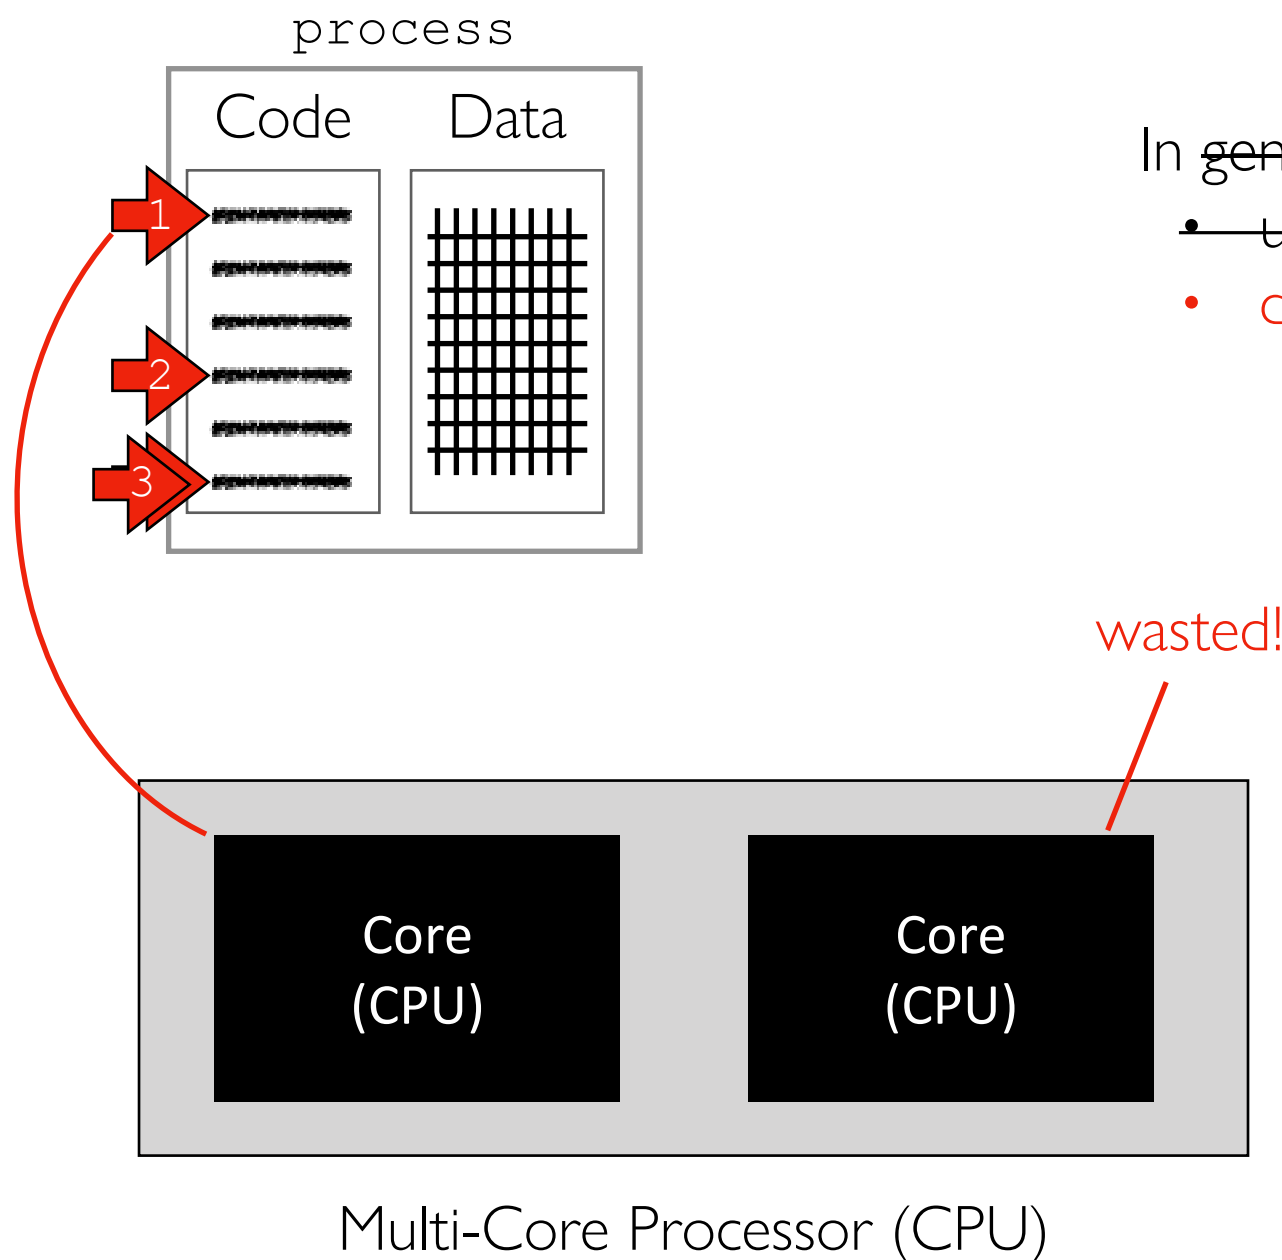


In general, threads help:

- use multiple cores
- do useful work when threads are blocking

Running: 1, 3  
Ready: 4  
Blocked: 2

# (1) Thread-level Parallelism



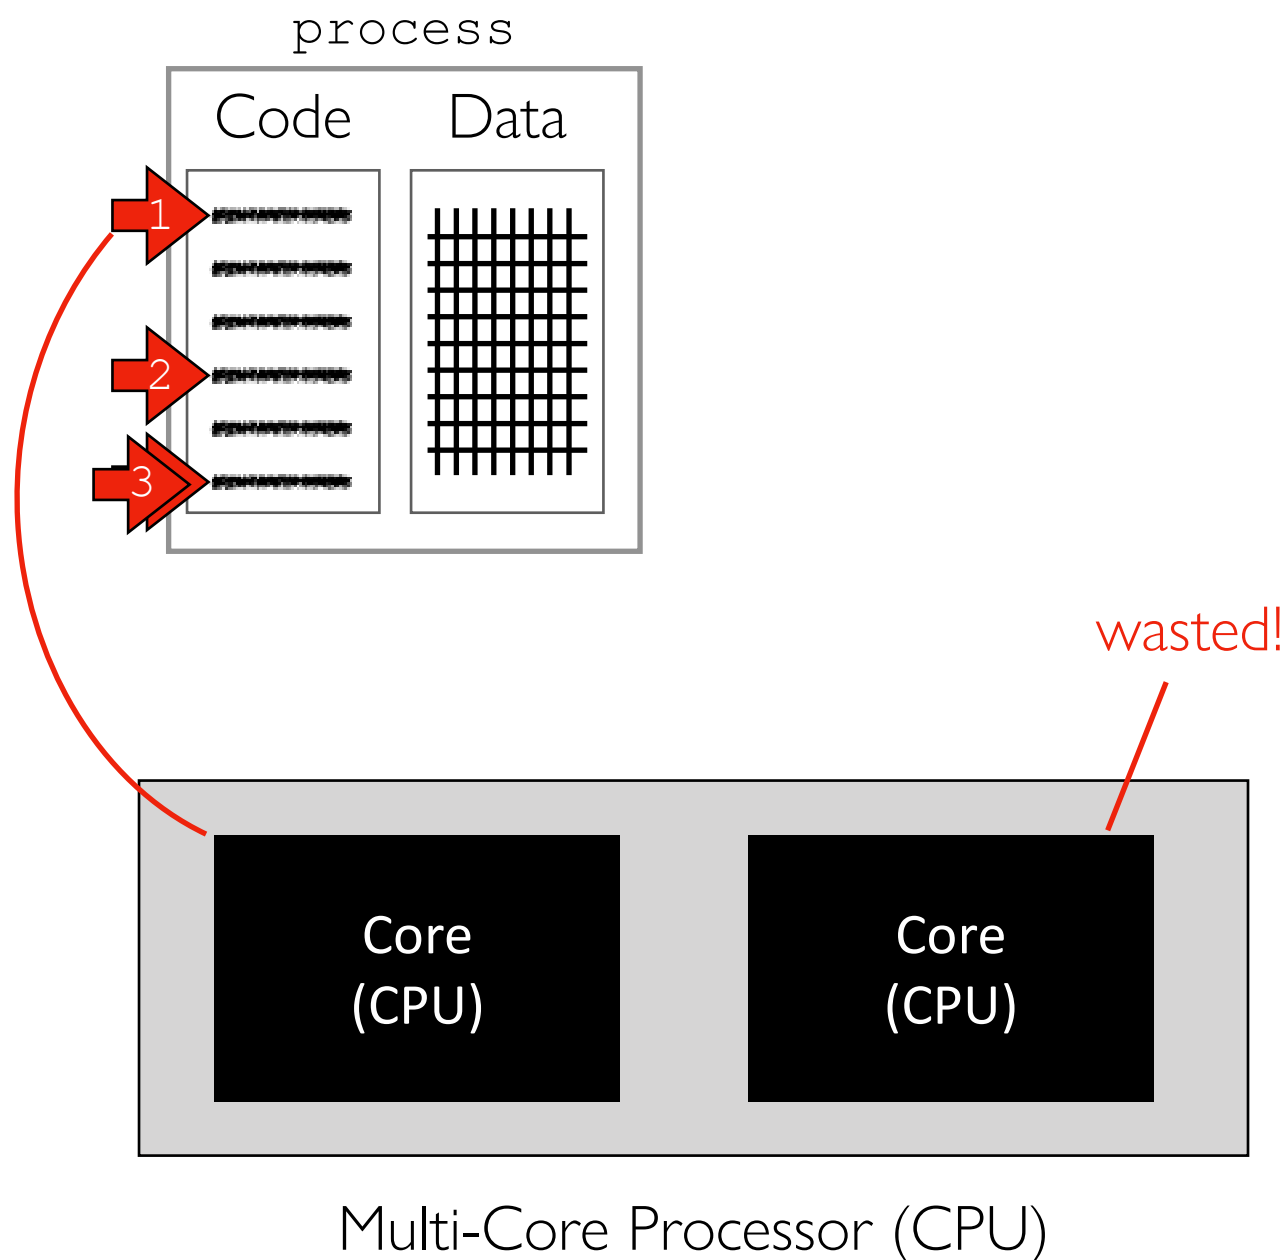
In ~~general~~ Python, threads help:

- ~~• use multiple cores~~
- do useful work when threads are blocking

Running: 1  
Ready: 3, 4  
Blocked: 2

# (1) Thread-level Parallelism

**recommendation:** don't use threads  
unless you learn a LOT more about  
multi-threading than covered in CS 320



Example: two countdown threads

```
import time
from threading import Thread
```

```
def f(name, n):
    for i in range(n):
        print(name, n-i)
        time.sleep(1)
```

```
# f("A", 3)
```

```
# f("B", 5)
```

```
t1 = Thread(target=f, args=("A", 3))
```

```
t2 = Thread(target=f, args=("B", 5))
```

```
t1.start()
```

```
t2.start()
```

```
t1.join()
```

```
t2.join()
```

Running: 1

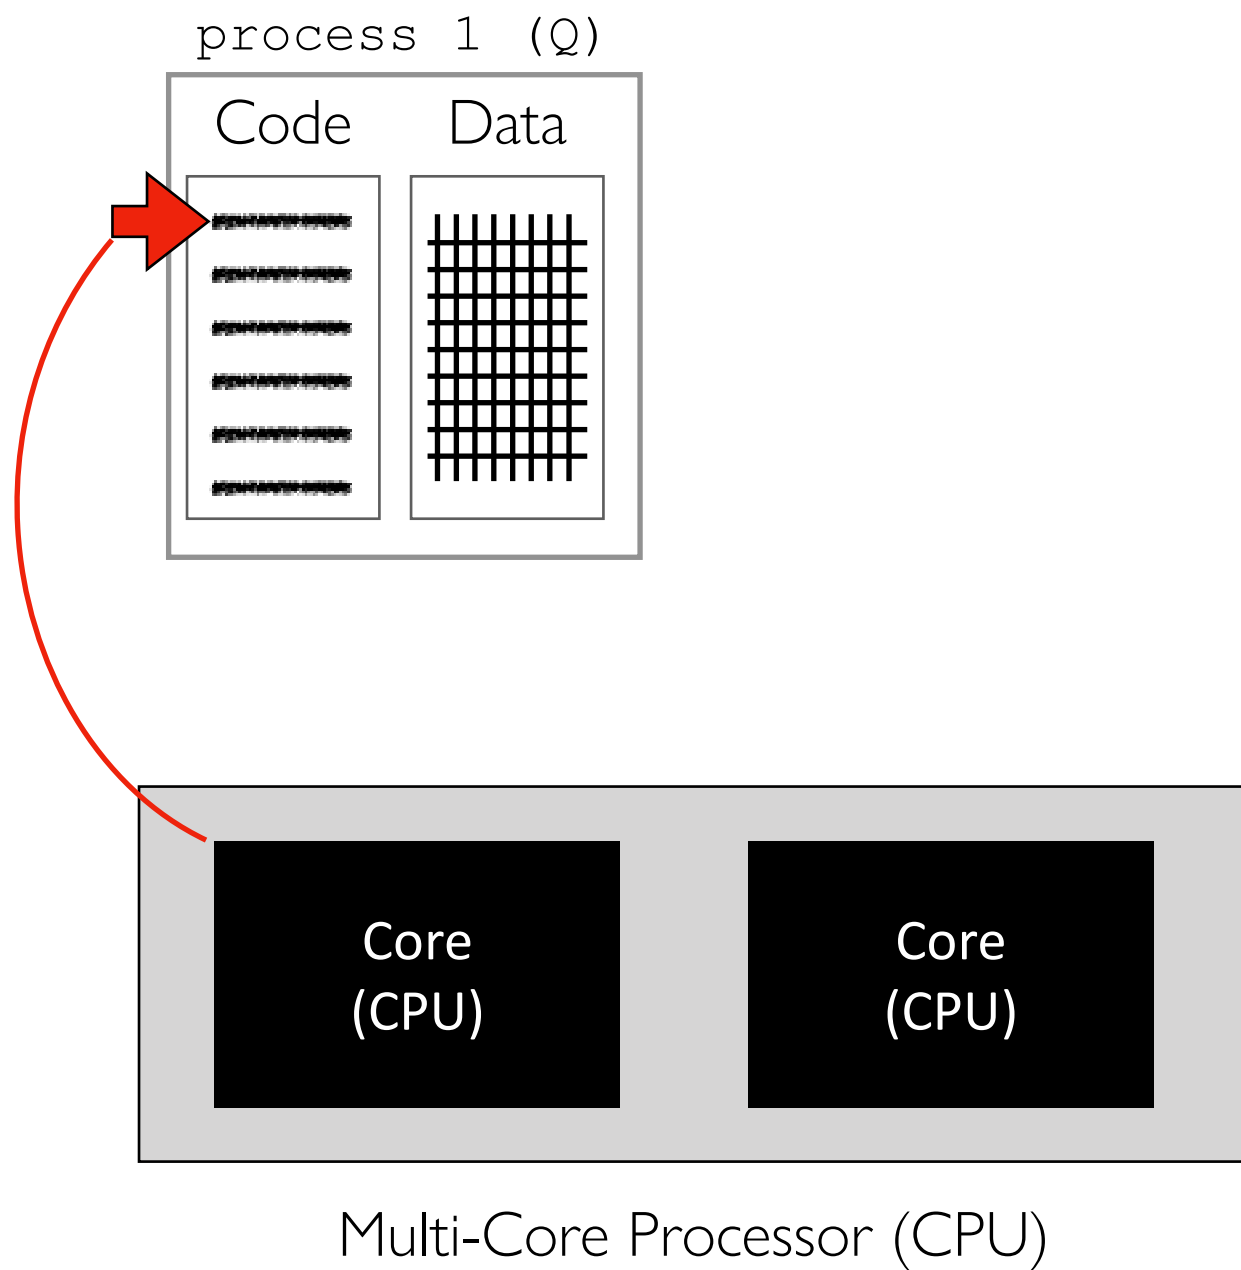
Ready: 3, 4

Blocked: 2

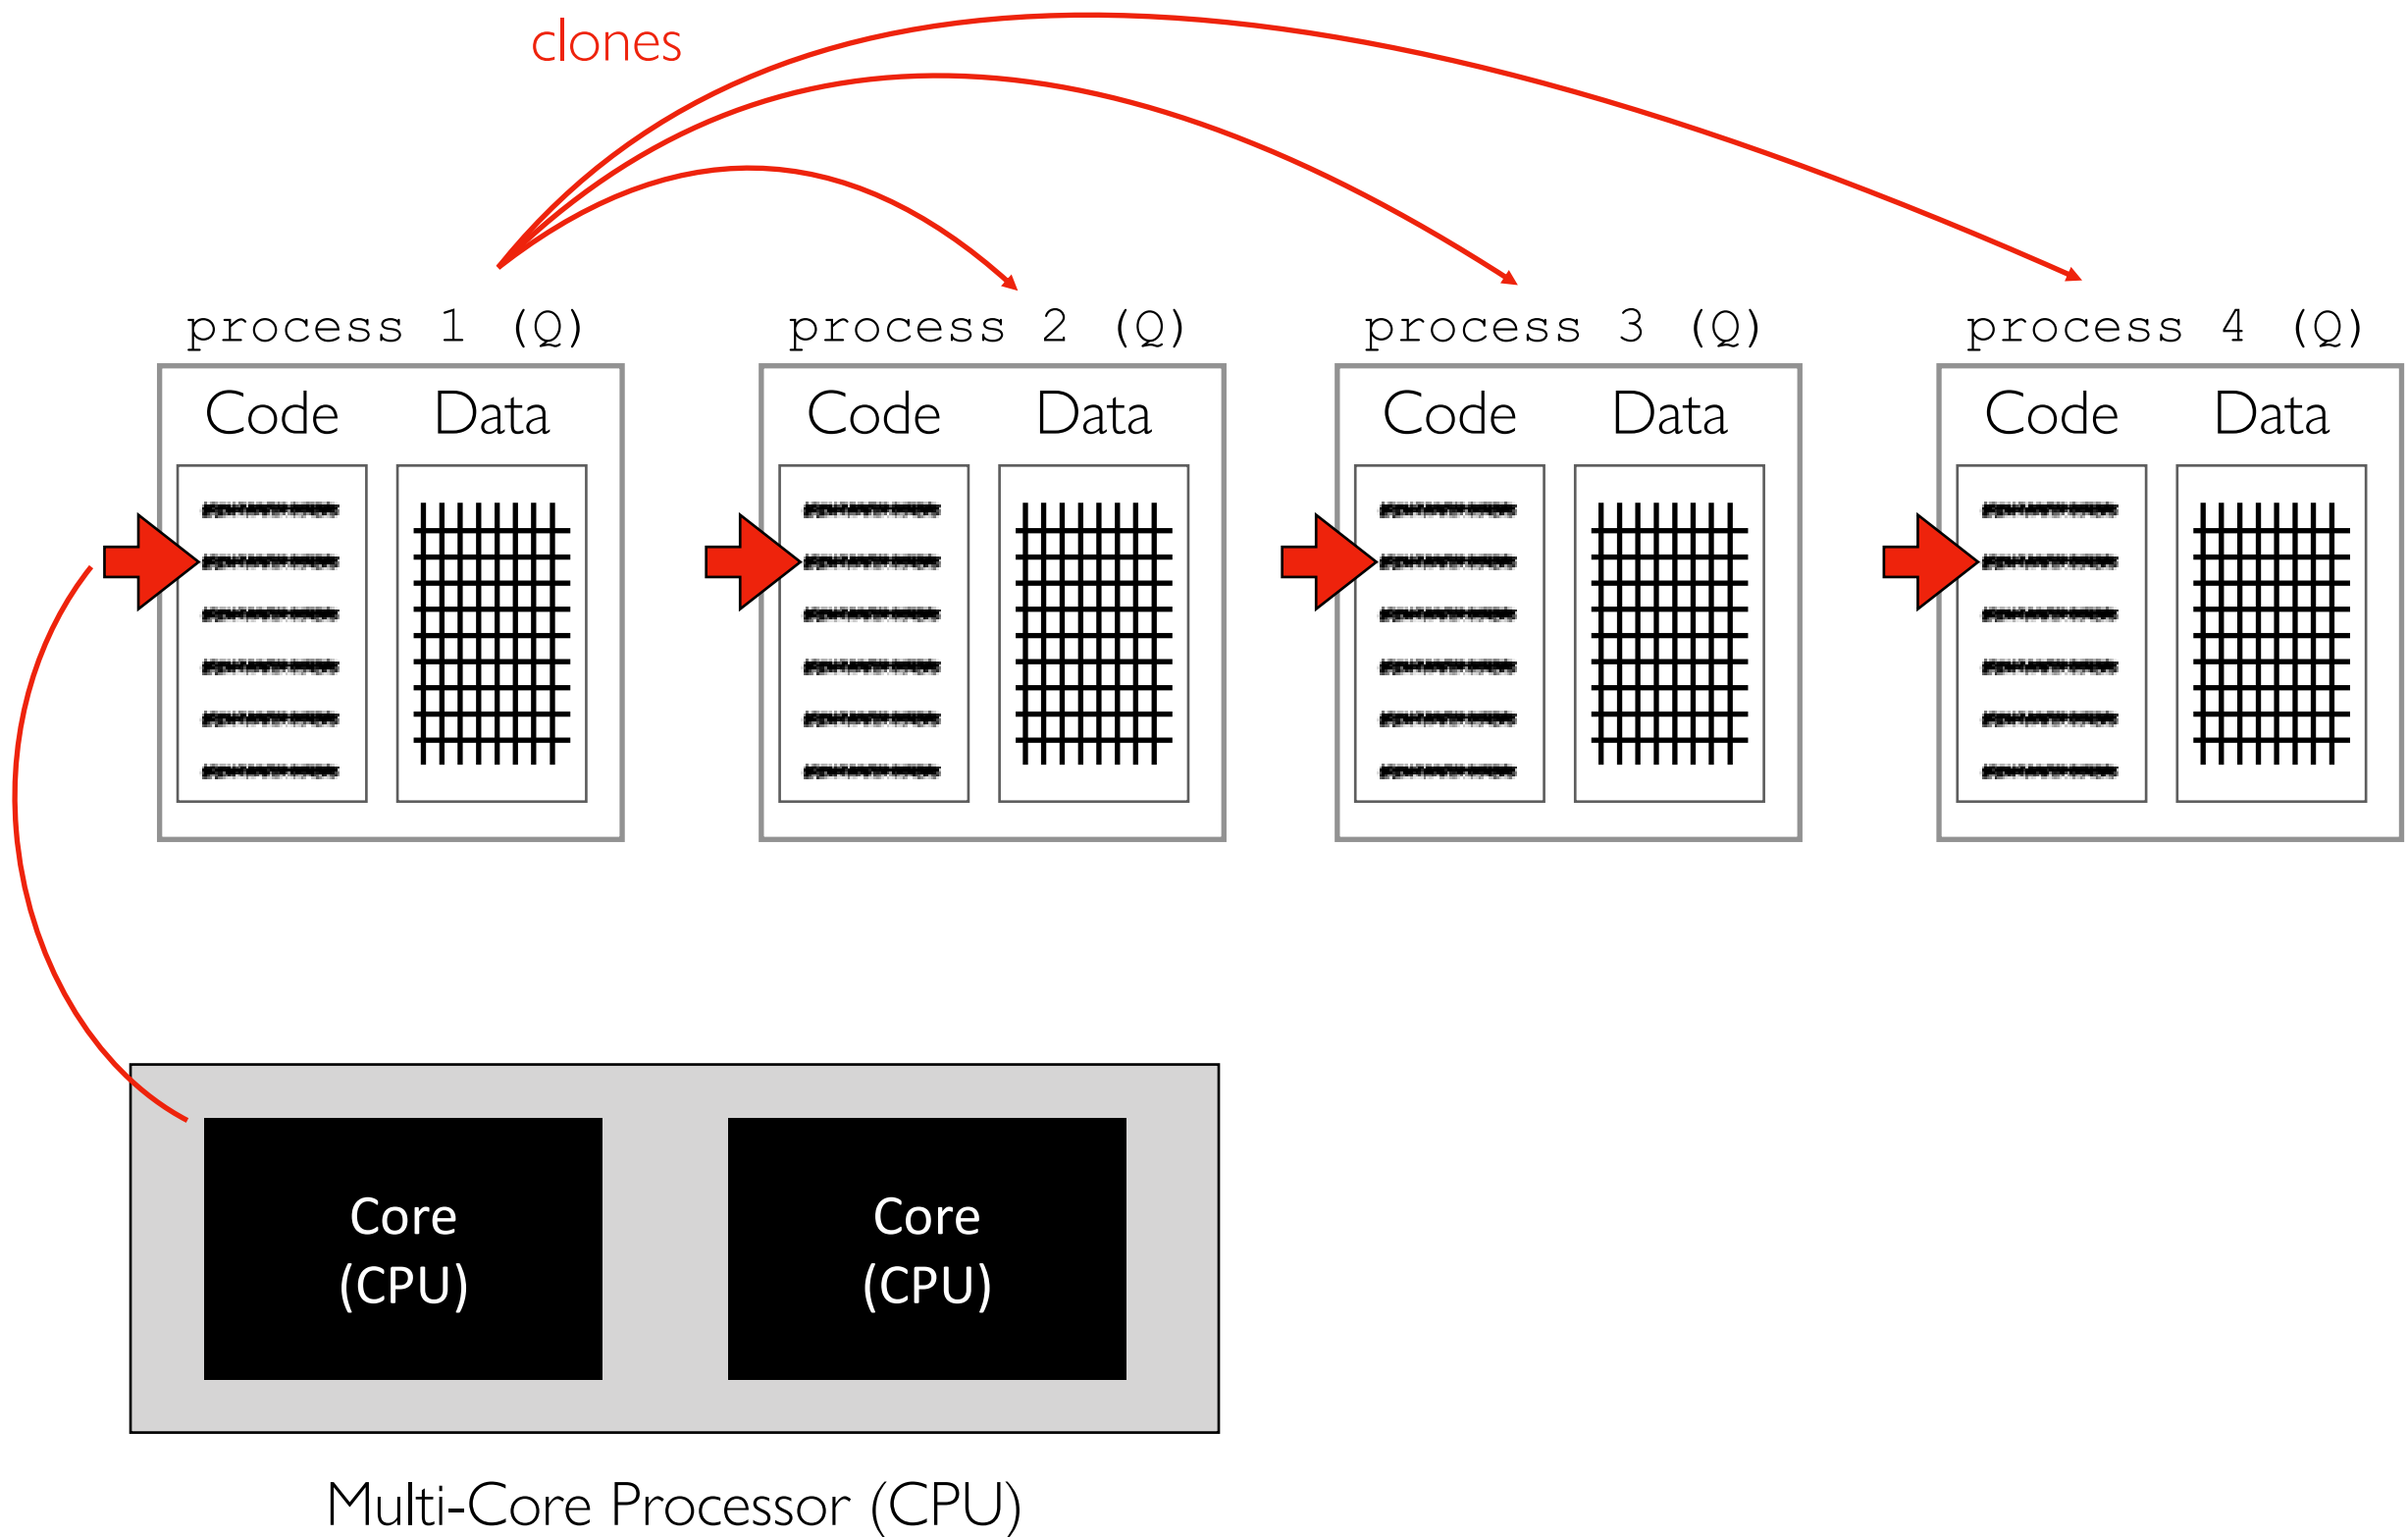
# Solution: Parallelism

- |   |                           |   |
|---|---------------------------|---|
| 1 | thread-level parallelism  | very complicated, not covered in detail |
| 2 | process-level parallelism | covered in CS 320                       |
| 3 | GPU parallelism           |   |

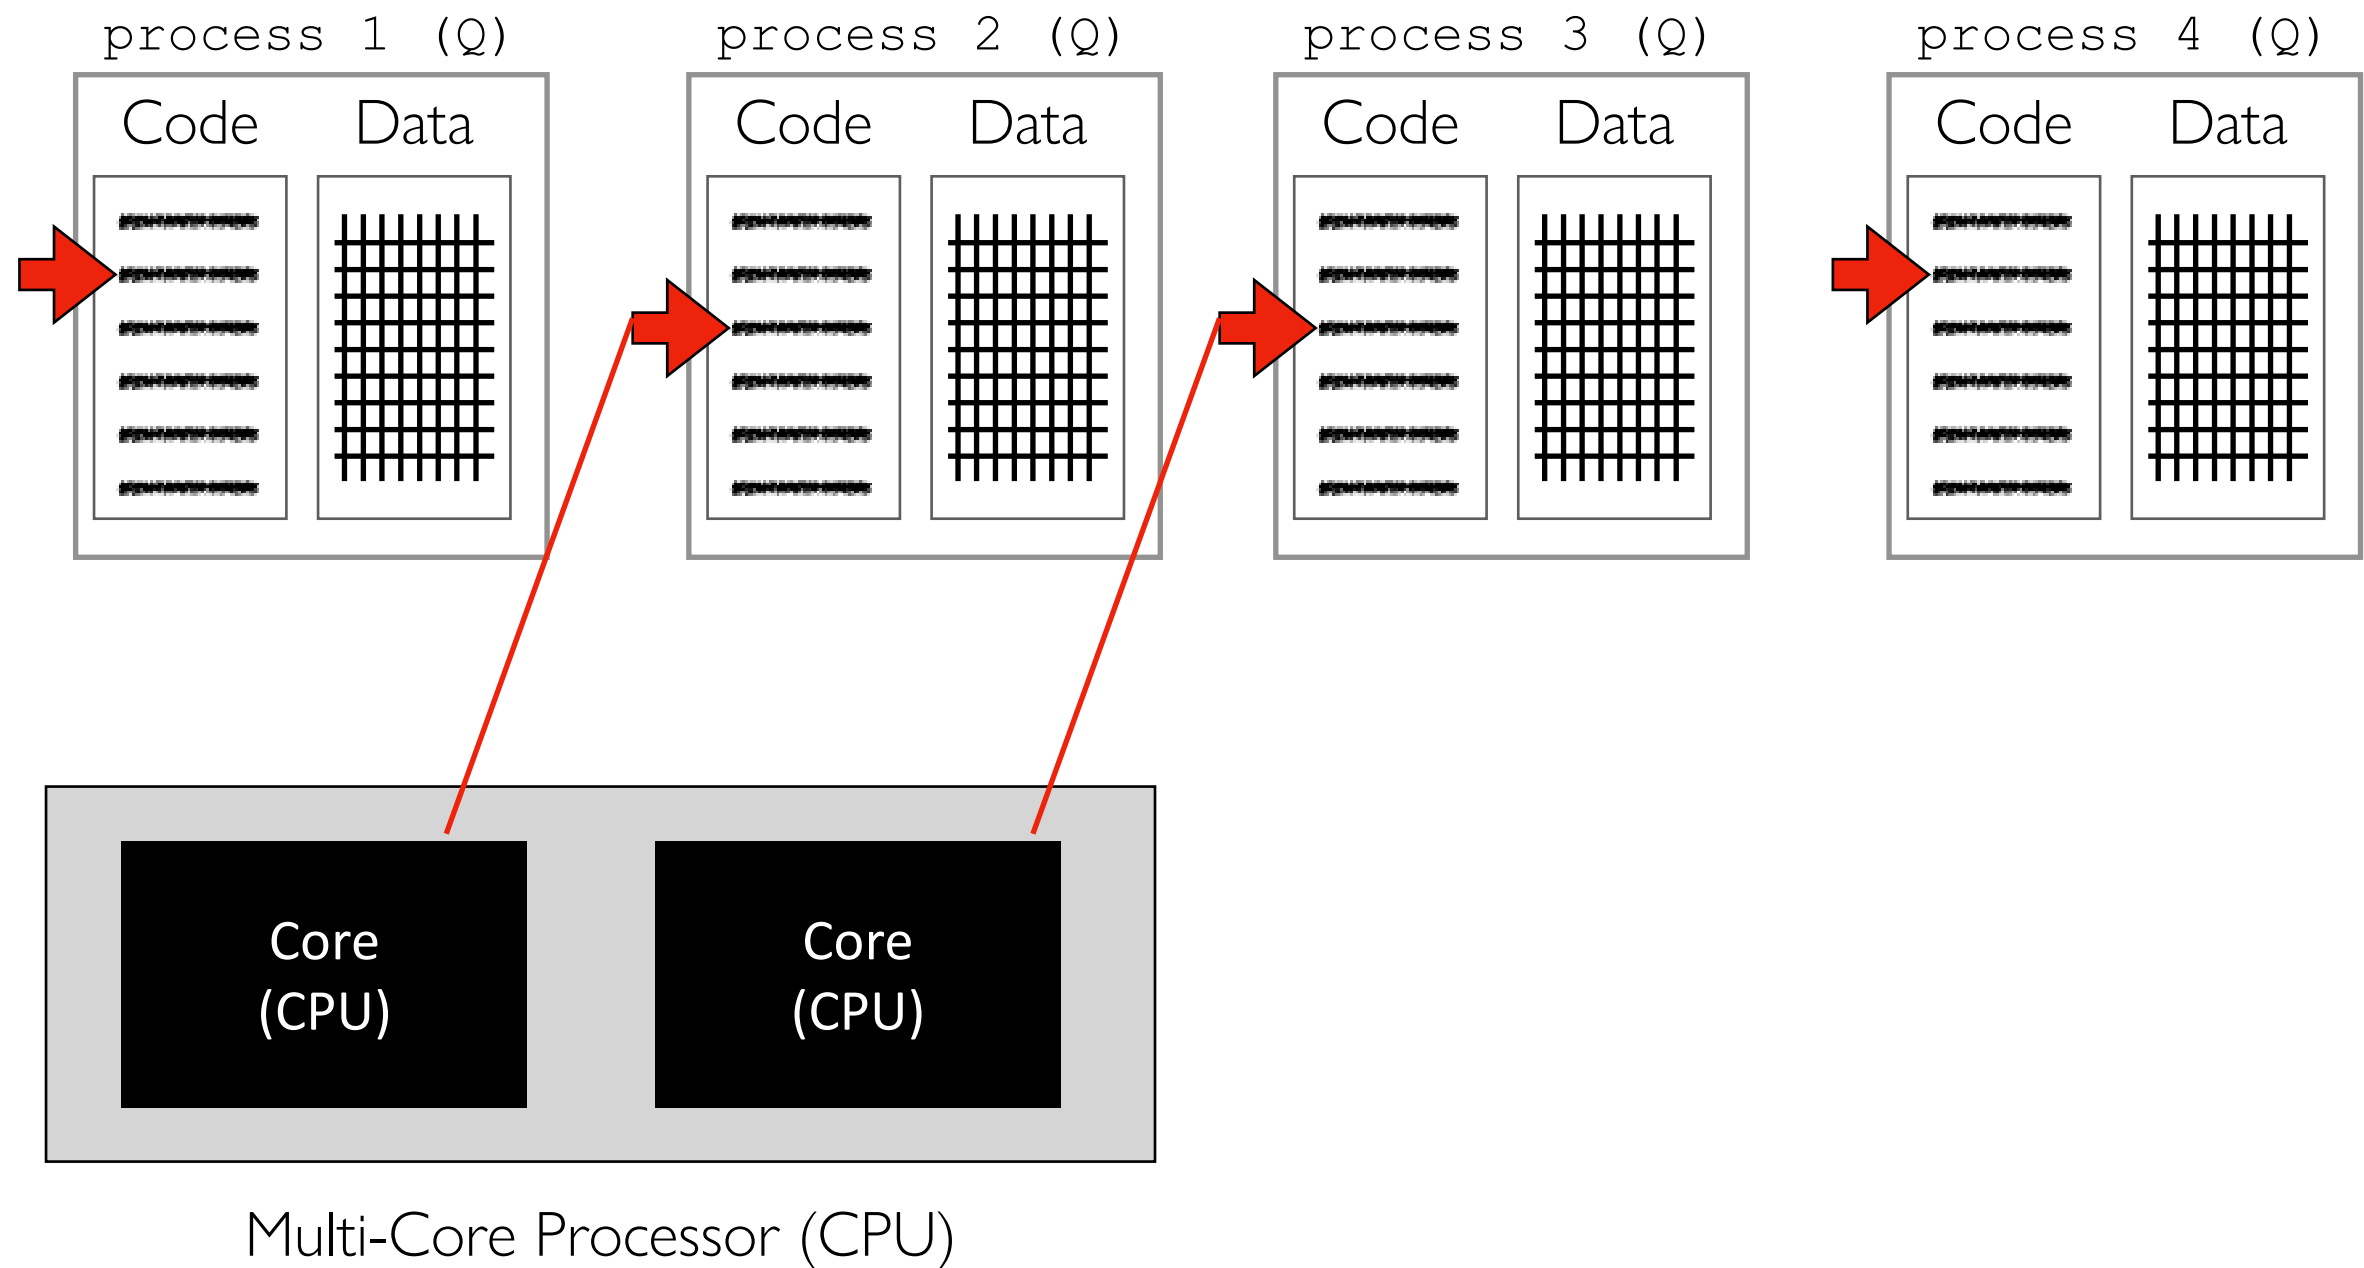
## (2) Process-level Parallelism



## (2) Process-level Parallelism

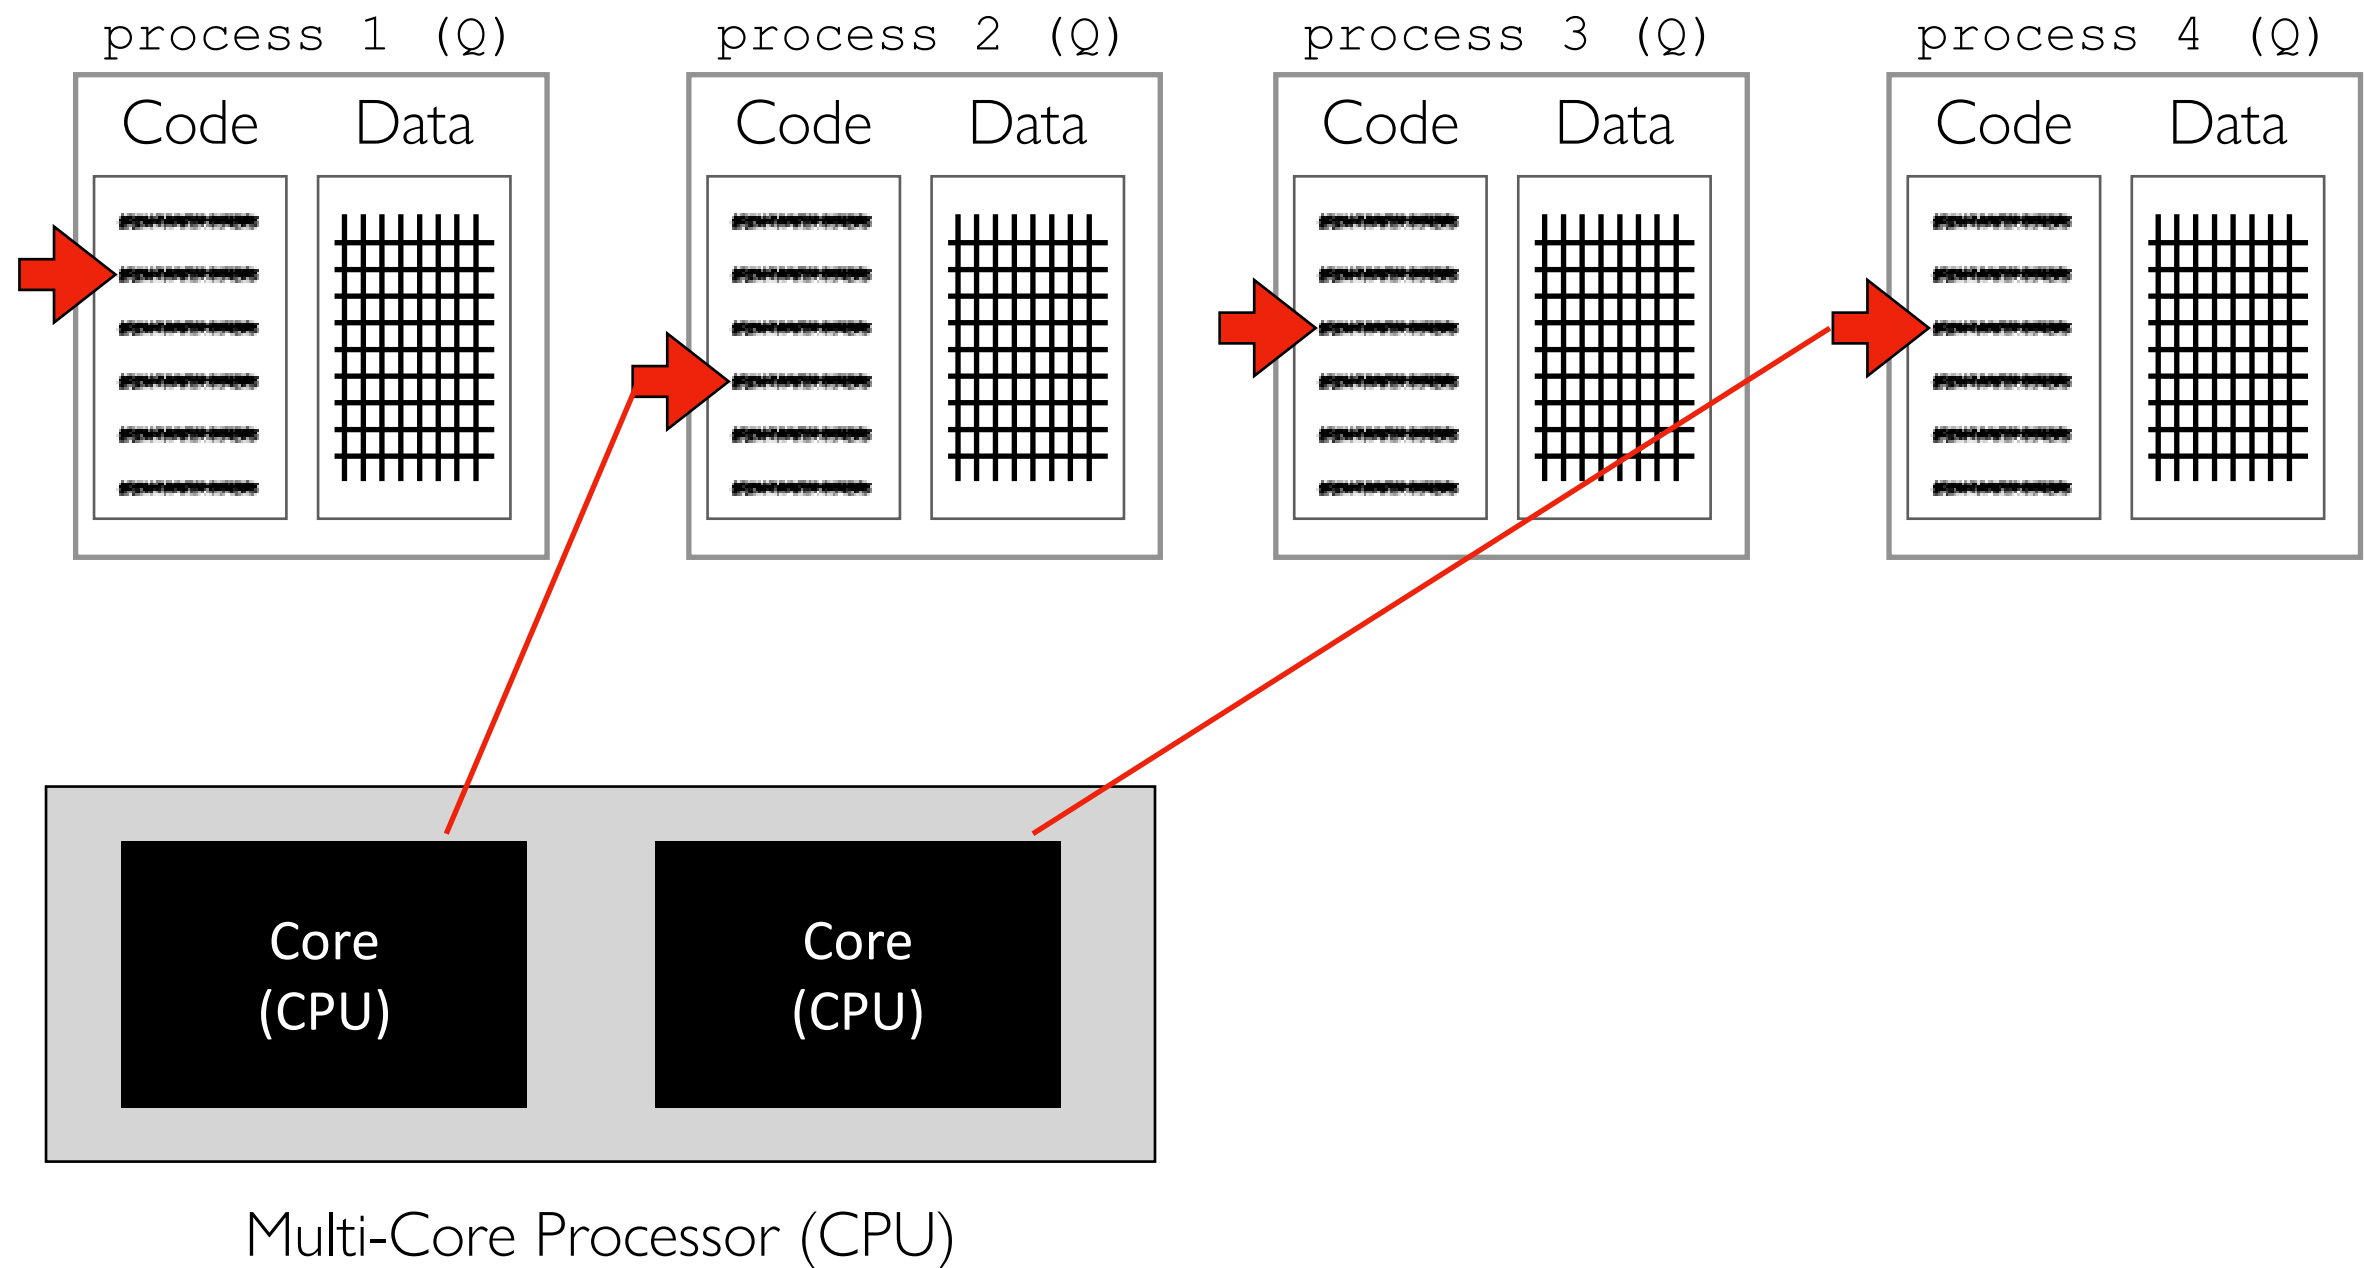


## (2) Process-level Parallelism

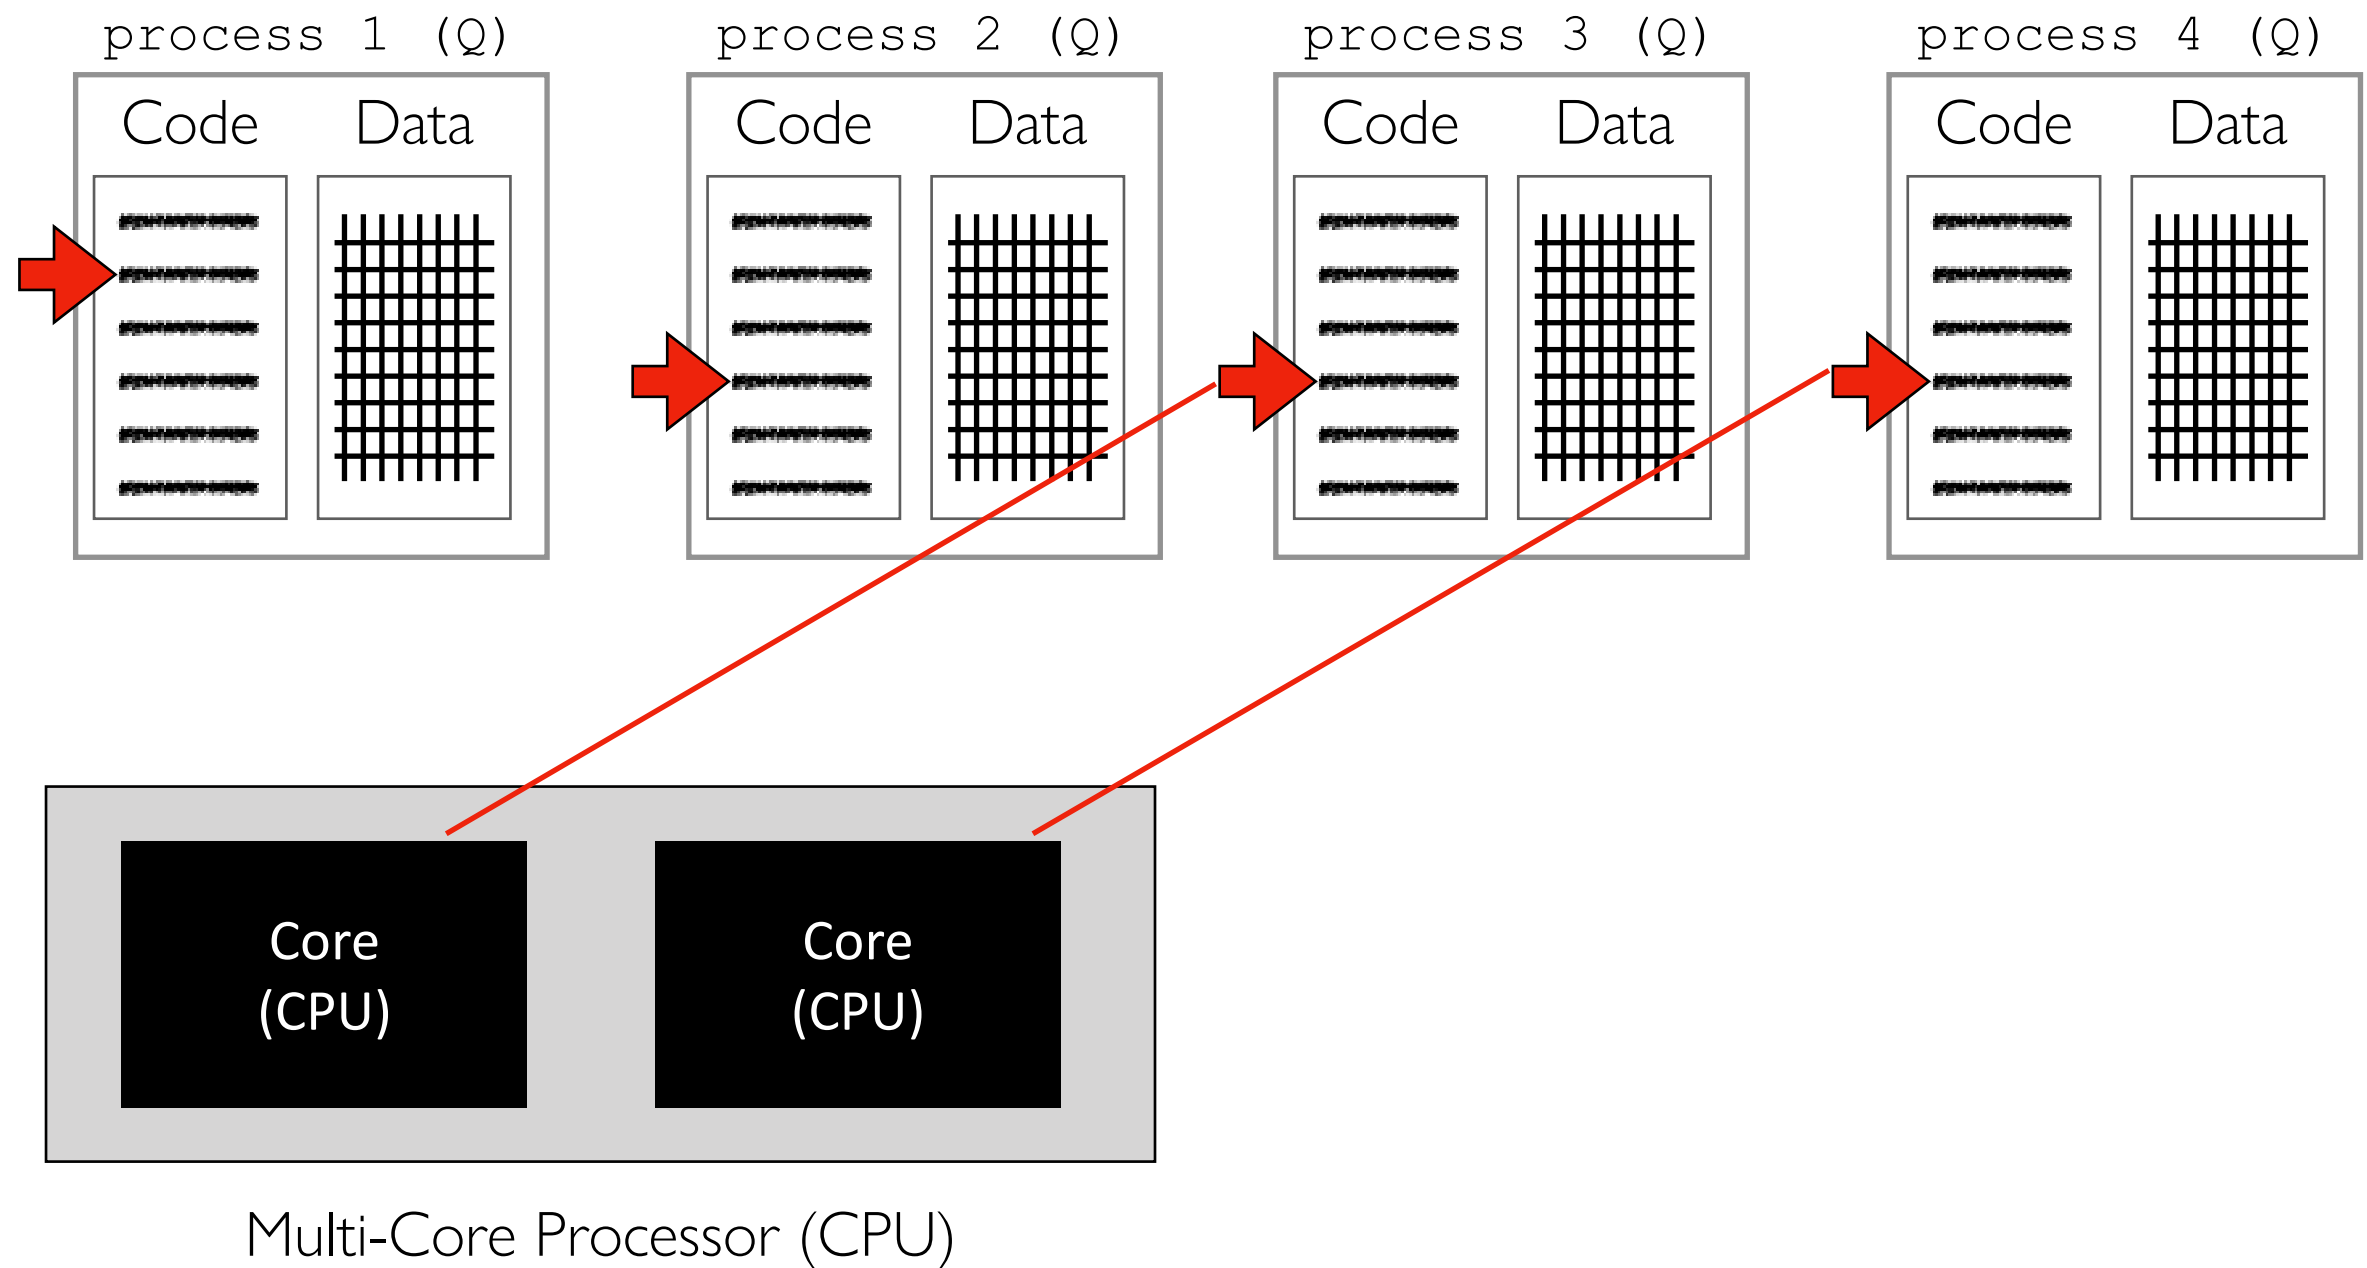




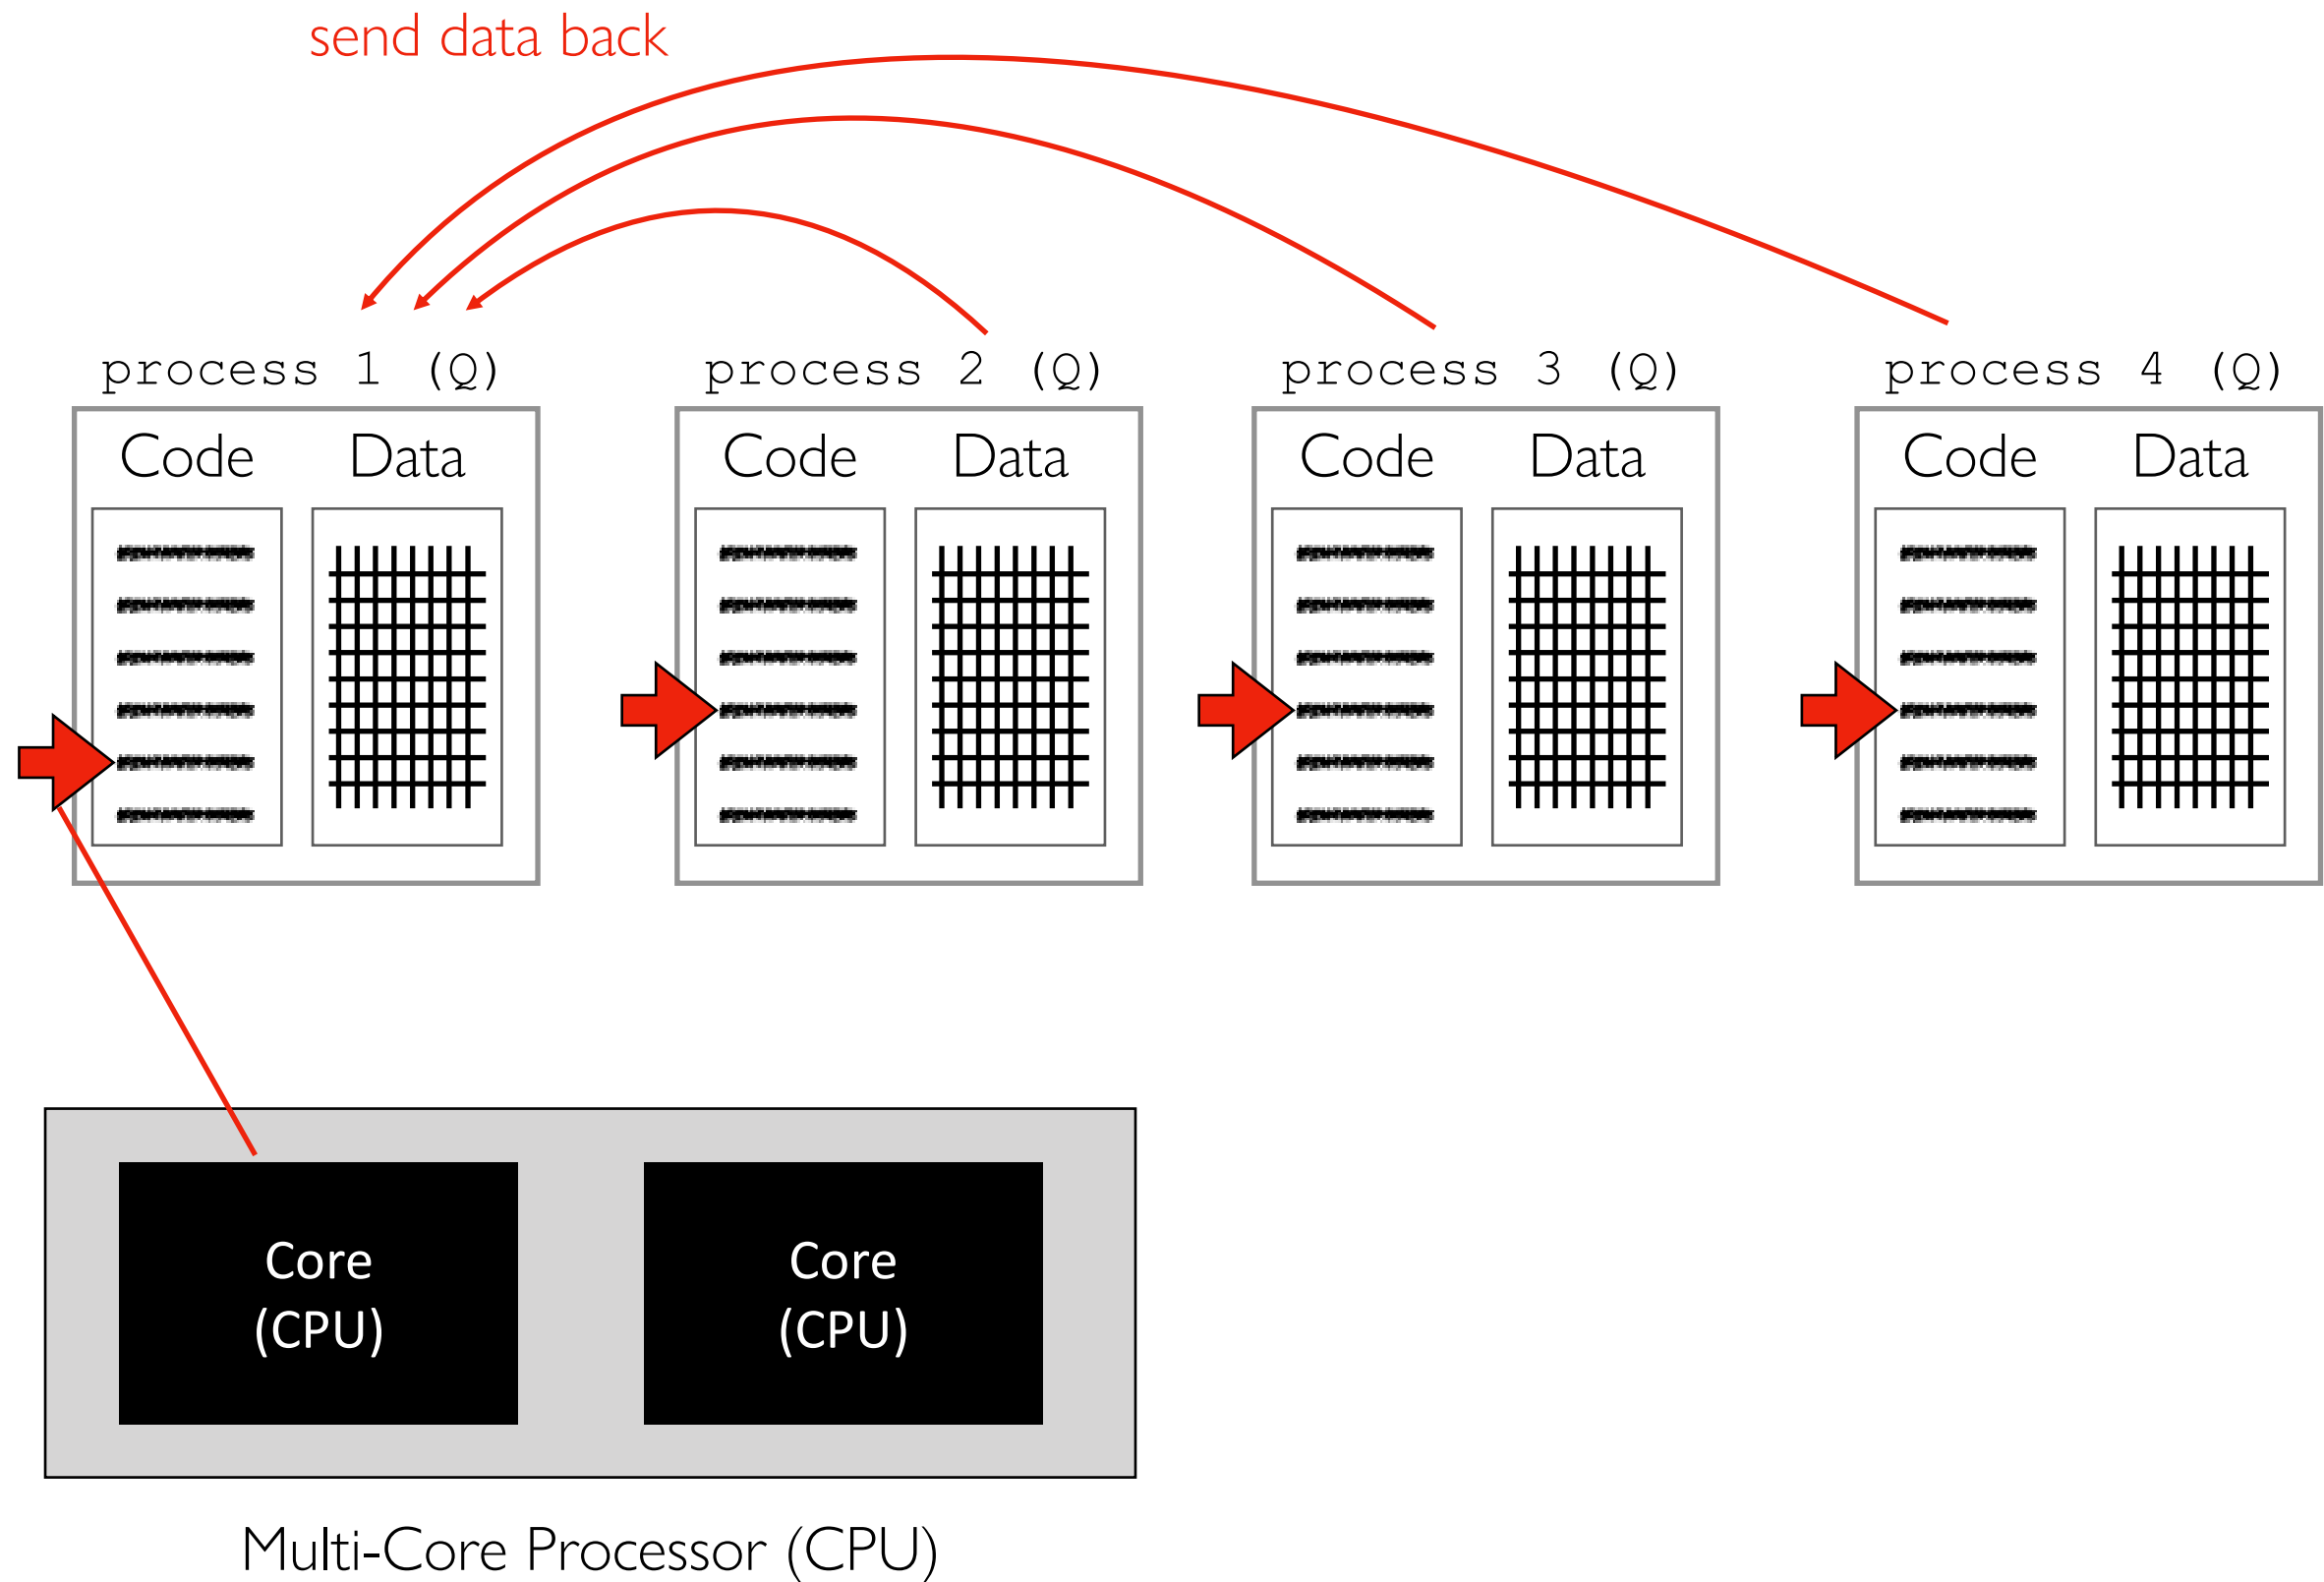
## (2) Process-level Parallelism




## (2) Process-level Parallelism



## (2) Process-level Parallelism




## (2) Process-level Parallelism



## (2) Process-level Parallelism


<https://docs.python.org/3/library/multiprocessing.html>



```
from multiprocessing import Pool

def f(x):
    return x*x

if __name__ == '__main__':
    with Pool(5) as p:
        print(p.map(f, [1, 2, 3]))
```




Multi-Core Processor (CPU)


# Solution: Parallelism

- |   |                           |   |
|---|---------------------------|---|
| 1 | thread-level parallelism  | very complicated, not covered in detail |
| 2 | process-level parallelism | covered in CS 320                       |
| 3 | GPU parallelism           |   |


### (3) GPU Parallelism



[https://en.wikipedia.org/wiki/Nvidia\\_Tesla](https://en.wikipedia.org/wiki/Nvidia_Tesla)



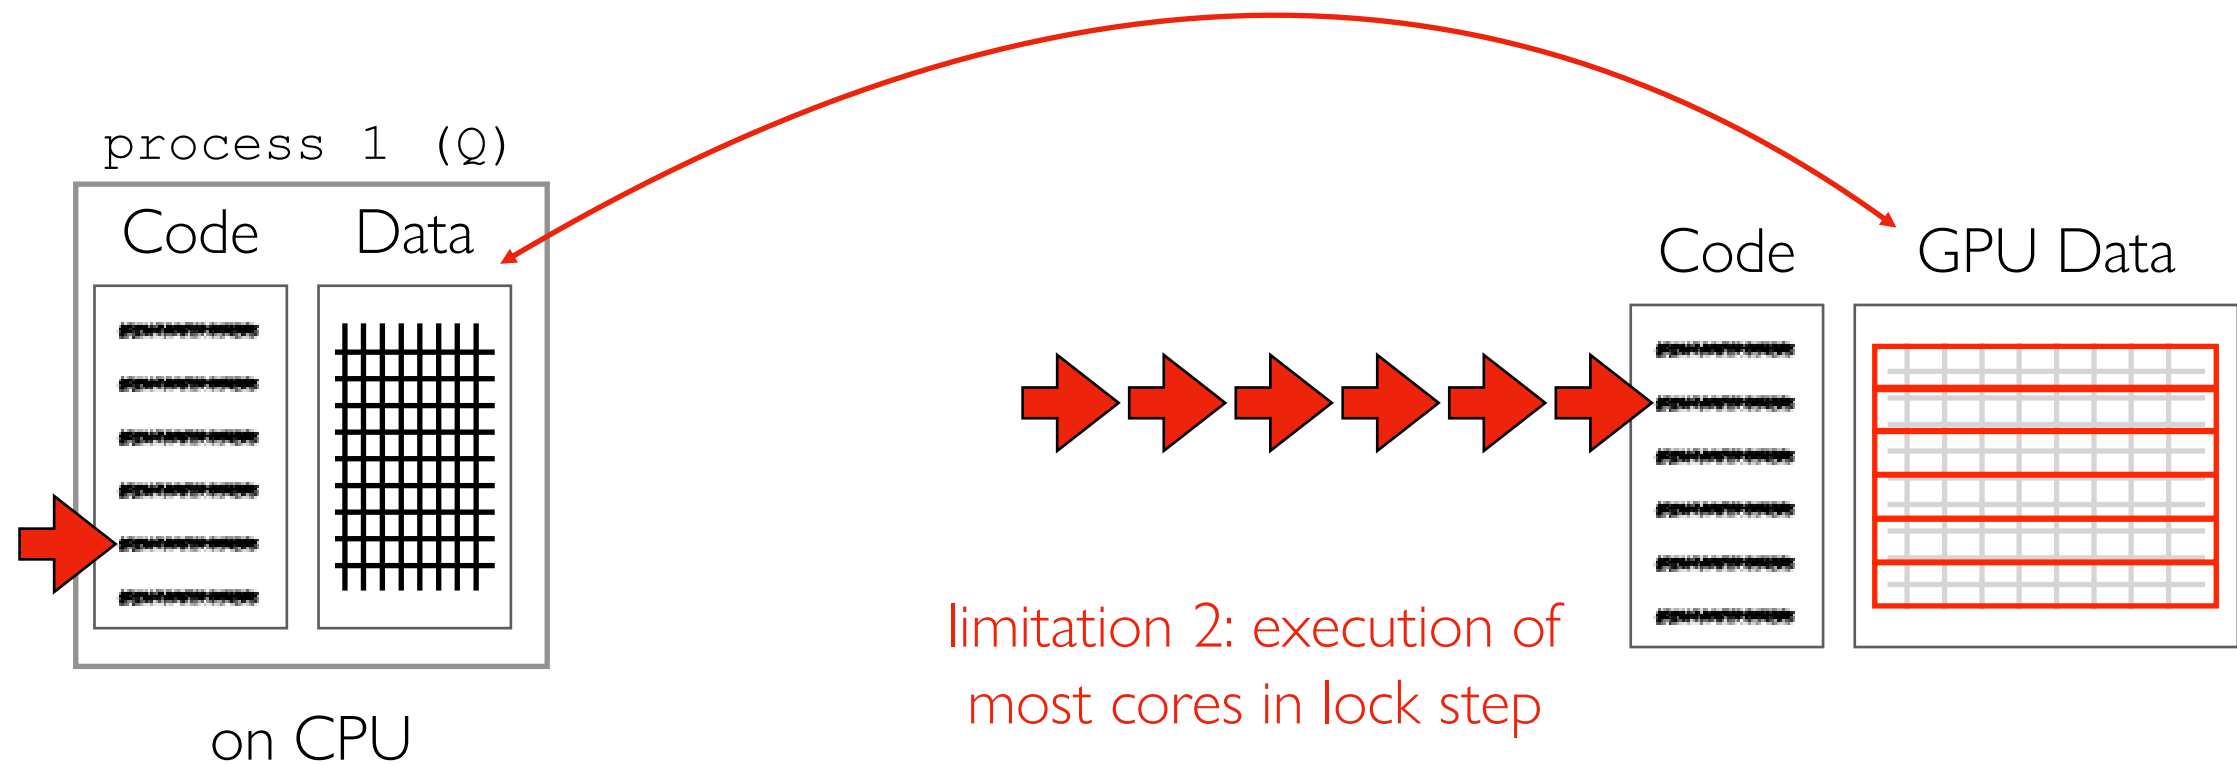
few cores that are fast,  
flexible, independent



many cores that are slow,  
float-optimized, coordinated

# GPU Limitations

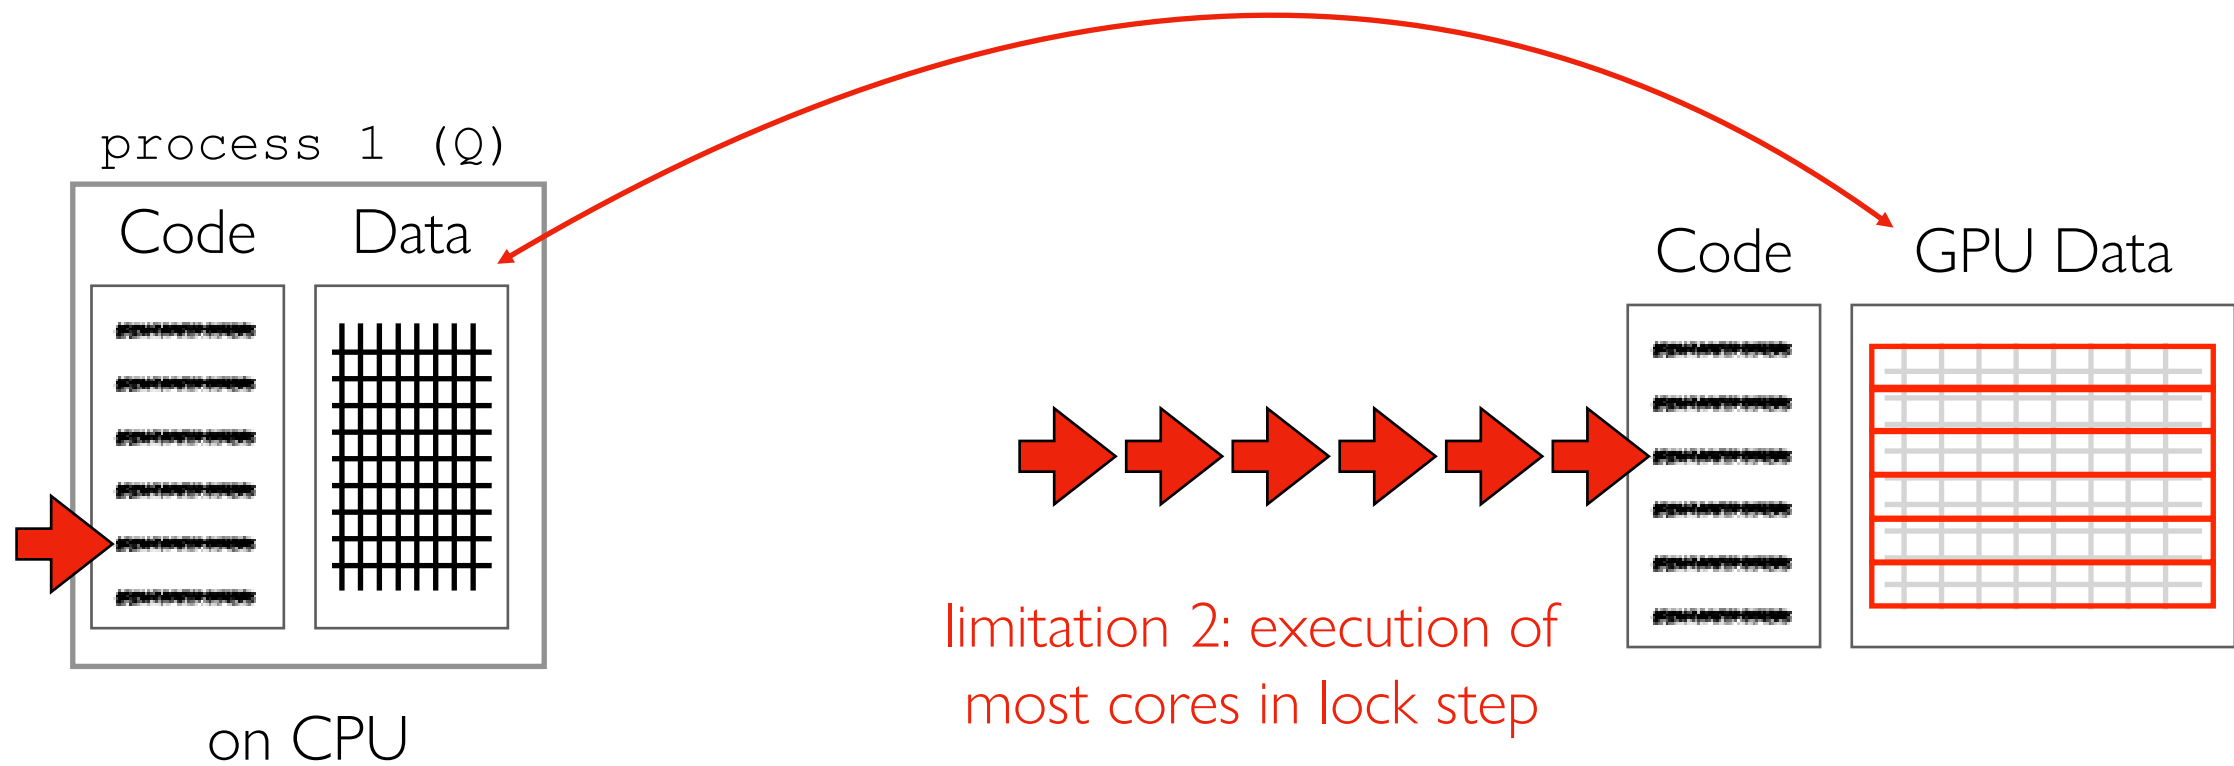
limitation 1: need to move data back and forth to GPU





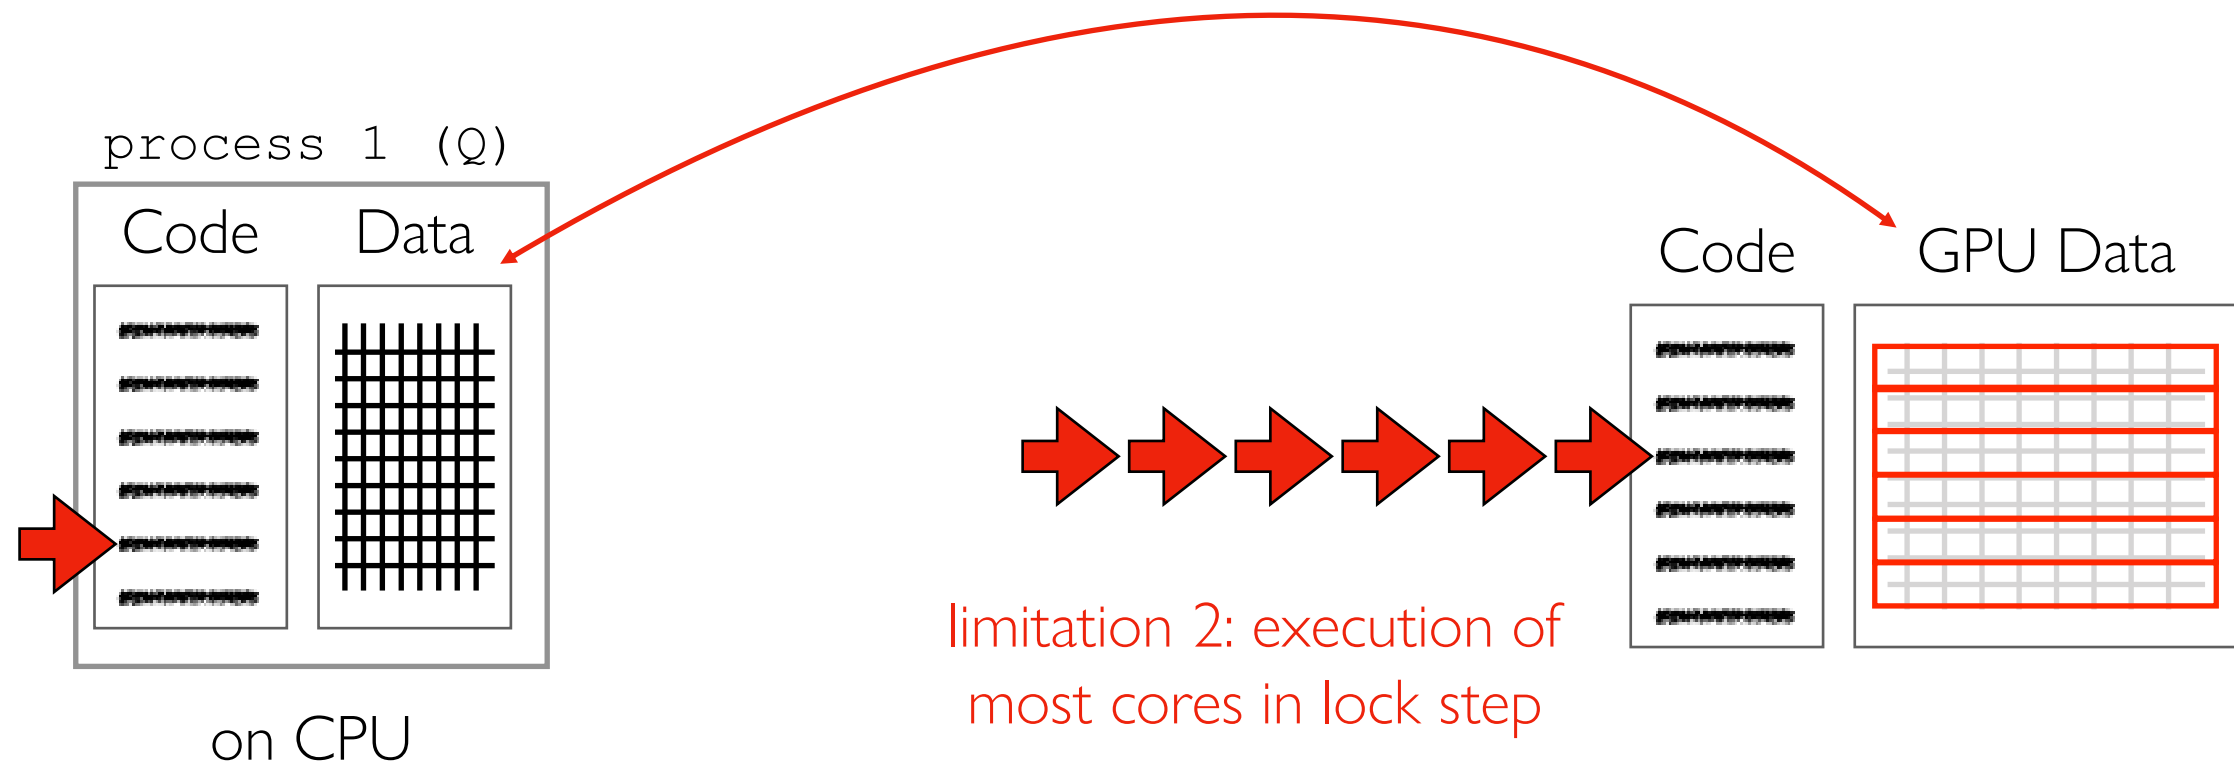
# GPU Limitations

limitation 1: need to move data back and forth to GPU



# GPU Limitations

limitation 1: need to move data back and forth to GPU




great use case:  
matrix multiplication

$$\begin{bmatrix} \text{row1} \\ \text{row2} \\ \dots \\ \text{rowN} \end{bmatrix} \cdot \begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} \text{output1} \\ \text{output2} \\ \dots \\ \text{outputN} \end{bmatrix}$$

multiply row 1 of matrix by vector,  
multiply row 2 of matrix by vector,  
multiply row 3 of matrix by vector,  
...

# GPU vs. CPU: Cost Comparison



Specifications	Intel® Core™ i7-6900K Processor Extreme Ed.	NVIDIA GeForce® GTX™ 1080 Ti
Base Clock Frequency	3.2 GHz	< 1.5 GHz
Cores	8	3584
Memory Bandwidth	64 GB/s	484 GB/s
Floating-Point Calculations	409 GFLOPS	11300 GFLOPS
Cost	~ \$1000.00	~ \$700.00

<https://sebastianraschka.com/books.html>

The GPU is 30% cheaper but 28x faster at floating-point operations!

# PyTorch

```
import numpy as np
import torch
A = np.random.normal(size=(1000,20))
x = np.random.normal(size=(20,1))
A = torch.from_numpy(A).to("cuda") # GPU
x = torch.from_numpy(x).to("cuda") # GPU
b = A @ x
b = b.to("cpu")
b
```

- CUDA: Compute Unified Device Architecture
- pytorch tensor is like numpy array
- .to("cuda") moves data to GPU
- .to("cpu") moves output back to CPU

# Parallelism

- 1 thread-level parallelism
- 2 process-level parallelism
- 3 GPU parallelism